



Semester 1 Examinations 2019/2020

Course Instance Code(s) Exam(s)	1CSD1, 1CSD2, 1SPE1, 1MAO2, 1MAI1 MSc in Computer Science (Data Analytics), MSc in Computer Science (Artificial Intelligence), MSc in Computer Science (Artificial Intelligence) - Online
Module Code(s) Module(s)	CT5120, CT5146 Introduction to Natural Language Processing, Introduction to Natural Language Processing - Online
Paper No.	1
Repeat Paper	No
External Examiner(s)	Professor Pier Luca Lanzi
Internal Examiner(s)	Dr. Michael Madden *Dr. Paul Buitelaar, Dr. John McCrae

Instructions: Answer all parts of all questions. There are 4 sections; each section is worth 25 marks (100 marks total). **Use a separate answer book for each section answered.**

Duration	2 hours
No. of Pages	5
Discipline(s)	Computer Science
Course Co-ordinator(s)	Dr. Enda Howley, Dr. Michael Schukat, Dr. James McDermott

Requirements:

Release in Exam Venue	Yes	<input checked="" type="checkbox"/>	No	<input type="checkbox"/>
MCQ	Yes	<input type="checkbox"/>	No	<input checked="" type="checkbox"/>
Handout	None			
Statistical/ Log Tables	None			
Cambridge Tables	None			
Graph Paper	None			
Log Graph Paper	None			
Other Materials	None	<input type="checkbox"/>		<input checked="" type="checkbox"/>
Graphic material in colour	Yes		No	

Introduction to Natural Language Processing

Exam Duration: 2 Hours

You must complete Sections 1 to 4

Section 1: Semantics

Instructions: Provide answers for questions 1A, 1B, 1C and 1D

Question 1A

10 Marks

Consider the following sentences:

The black horse jumped over the white fence.

The black fence was too high for the black horse.

The white horse could not have jumped over the black fence.

The black horse jumped successfully over the black fence.

The white fence was also too high for the white horse.

Using cosine similarity, compute the distributional semantic distance between the vectors for 'black' and 'white' constructed on the basis of the sentences above with a context window of two tokens. Give the vectors as well as the similarity score.

Question 1B

5 Marks

How can dependency parsing be used in a vector space model?

~~Question 1C~~

5 Marks

~~What is the opposite of a hyponym?~~

Question 1D

5 Marks

Give 2 examples of types of semantic roles and construct a sentence in which both of these can be used. Annotate the relevant words in the sentence for each of the 2 roles.

PTO

Section 2: Part-of-speech tagging

Instructions: Provide answers for question 2A, 2B and 2C

Question 2A

10 Marks

Consider a Hidden Markov Model with the following probabilities (Start designates the start state):

$p(w_i t_i)$	$w_i = \text{tag}$	$w_i = \text{this}$	$w_i = \text{text}$	$w_i = \text{herring}$
$t_i = N$	0.2	0.2	0.5	0.1
$t_i = V$	0.6	0.1	0.1	0.2
$t_i = O$	0.1	0.7	0.1	0.1

$p(t_i t_{i-1})$	$t_{i-1} = N$	$t_{i-1} = V$	$t_{i-1} = O$	$t_{i-1} = \text{Start}$
$t_i = N$	0.4	0.3	0.4	0.4
$t_i = V$	0.2	0.2	0.3	0.5
$t_i = O$	0.4	0.5	0.3	0.1

By using the Viterbi algorithm or otherwise, what is the most likely sequence of tags for the text “tag this text”?

Question 2B

5 Marks

What is the advantage of the Viterbi algorithm over an exhaustive search of all possible part-of-speech tag combinations?

Question 2C

10 Marks

Give one advantage and disadvantage of unsupervised learning of a Hidden Markov Model by means of the Baum-Welch algorithm as opposed to supervised learning of the model. Why might you wish to combine supervised and unsupervised learning for Hidden Markov Models?

PTO

Section 3: Sentiment Analysis

Instructions: Provide answers for question 3A, 3B and 3C

Consider the following sentences from product reviews with sentiment scores, using a range of -1, 0, +1:

- 1 *The food was great, but the service in this restaurant was unfriendly.*
- +1 *The steak was cooked to perfection and the service was great.*
- 0 *The best part of the meal was the beer.*

Question 3A

10 Marks

Create a sentiment lexicon on the basis of the review sentences given above.

Question 3B

10 Marks

Recall that a count vector represents the proportions of negative and positive sentiment in a review text. Using the sentiment lexicon you created in question 3A, calculate the count vector for each review sentence given above.

Question 3C

5 Marks

Give an example of *aspect-based* sentiment analysis from the review sentences given above.

PTO

Section 4: Information Extraction & Knowledge Graphs

Instructions: Provide answers for questions 4A, 4B and 4C

Consider the following sentence:

David Robert Joseph Beckham is a former professional footballer who played for Manchester United and the England national team.

Question 4A

10 Marks

Annotate the sentence above for Named Entity types 'person' (PER), 'organization' (ORG) and 'location' (LOC) by the use of the IOB tagging scheme. Explain the reasoning behind your annotations.

Question 4B

10 Marks

The sentence above provides a positive instance for extracting the 'played for' relation between a 'person' (footballer) and an 'organization' (football team). Give a negative instance for this relation.

Question 4C

5 Marks

How can clustering be used in taxonomy extraction? What are potential problems with this approach?

END