Ollscoil na Gaillimhe
University of Galway

Advanced Topics in NLP

Introduction

Dr. Paul Buitelaar & Dr. Omnia Zayed
Data Science Institute
University of Galway

University
*of*Galway.ie

Learning Objectives of this Course

- Gain insights into knowledge extraction from text, in particular around entities and relations
- Gain insights into opinion mining, in particular on emotion analysis, dynamic identification of aspect and analysis of figurative language
- Gain insights into language generation, in particular in the context of machine translation and chatbot development

Ollscoil na Gaillimhe
University of Galway

Ollscoil na Gaillimhe
University of Galway

Administrative Issues

Ollscoil na Gaillimhe
University of Galway

Lecturers

Dr. Paul Buitelaar & Dr. Omnia Zayed

Lecture Plan

Date		Lecture
9/1/23		Introduction
16/1/23		Knowledge Extraction I: Entities
23/1/23		Knowledge Extraction II: Relations
30/1/23		Opinion Mining I: Emotions
6/2/23		bank holiday
13/2/23		Opinion Mining II: Aspect
20/2/23		Opinion Mining III: Figurative Language
27/2/23		Language Generation I: Machine Translation I
6/3/23		Language Generation II: Machine Translation II + NLG
13/3/23		Language Generation III: Dialog Systems
20/3/23		Summary
27/3/23		Industry Talk

Ollscoil na Gaillimhe
University of Galway

Labs

Time & Venue: Fridays 11am-1pm, IT102

Practical exercises covering the course content

Prerequisites are basic knowledge of Python, see:
https://docs.python.org/3/tutorial/

We will use Google colab Jupyter notebook, see:
https://colab.research.google.com

Dhairya Dalal

Ali Hatami

Lab Exercises

Date		Lab
13/1/23		intro to base NLP tools and methods
20/1/23		Named Entity Recognition
27/1/23		relation prediction
3/2/23		emotion classification
10/2/23		no lab - bank holiday
17/2/23		aspect-based sentiment analysis
24/2/23		metaphor classification
3/3/23		build a machine translation model step by step
10/3/23		machine translation evaluation
17/3/23		no lab - bank holiday
24/3/23		chatbot development with RASA

Ollscoil na Gaillimhe
University of Galway

Assignments

Two assignments

- Assignment 1: Released Jan 30th, Due Feb 20th
- Assignment 2: Released Mar 6th, Due Apr 3rd

Assignments count for 50% of final grade

Ollscoil na Gaillimhe
University of Galway

Recommended Reading

Lectures

- Jurafsky and Martin, ***SPEECH and LANGUAGE PROCESSING: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition***, 3rd edition: https://web.stanford.edu/~jurafsky/slp3/ed3book.pdf
- Chris Manning and Hinrich Schütze, ***Foundations of Statistical Natural Language Processing***, MIT Press. Cambridge, MA: May 1999: https://nlp.stanford.edu/fsnlp/

Labs

- ***Dive into Deep Learning***: https://d2l.ai/

Ollscoil na Gaillimhe
University of Galway

Summary of Intro NLP

“What we learned so far”

Ollscoil na Gaillimhe
University of Galway

Buzz Groups

What have we learned in the Intro NLP course?

Natural Language Processing

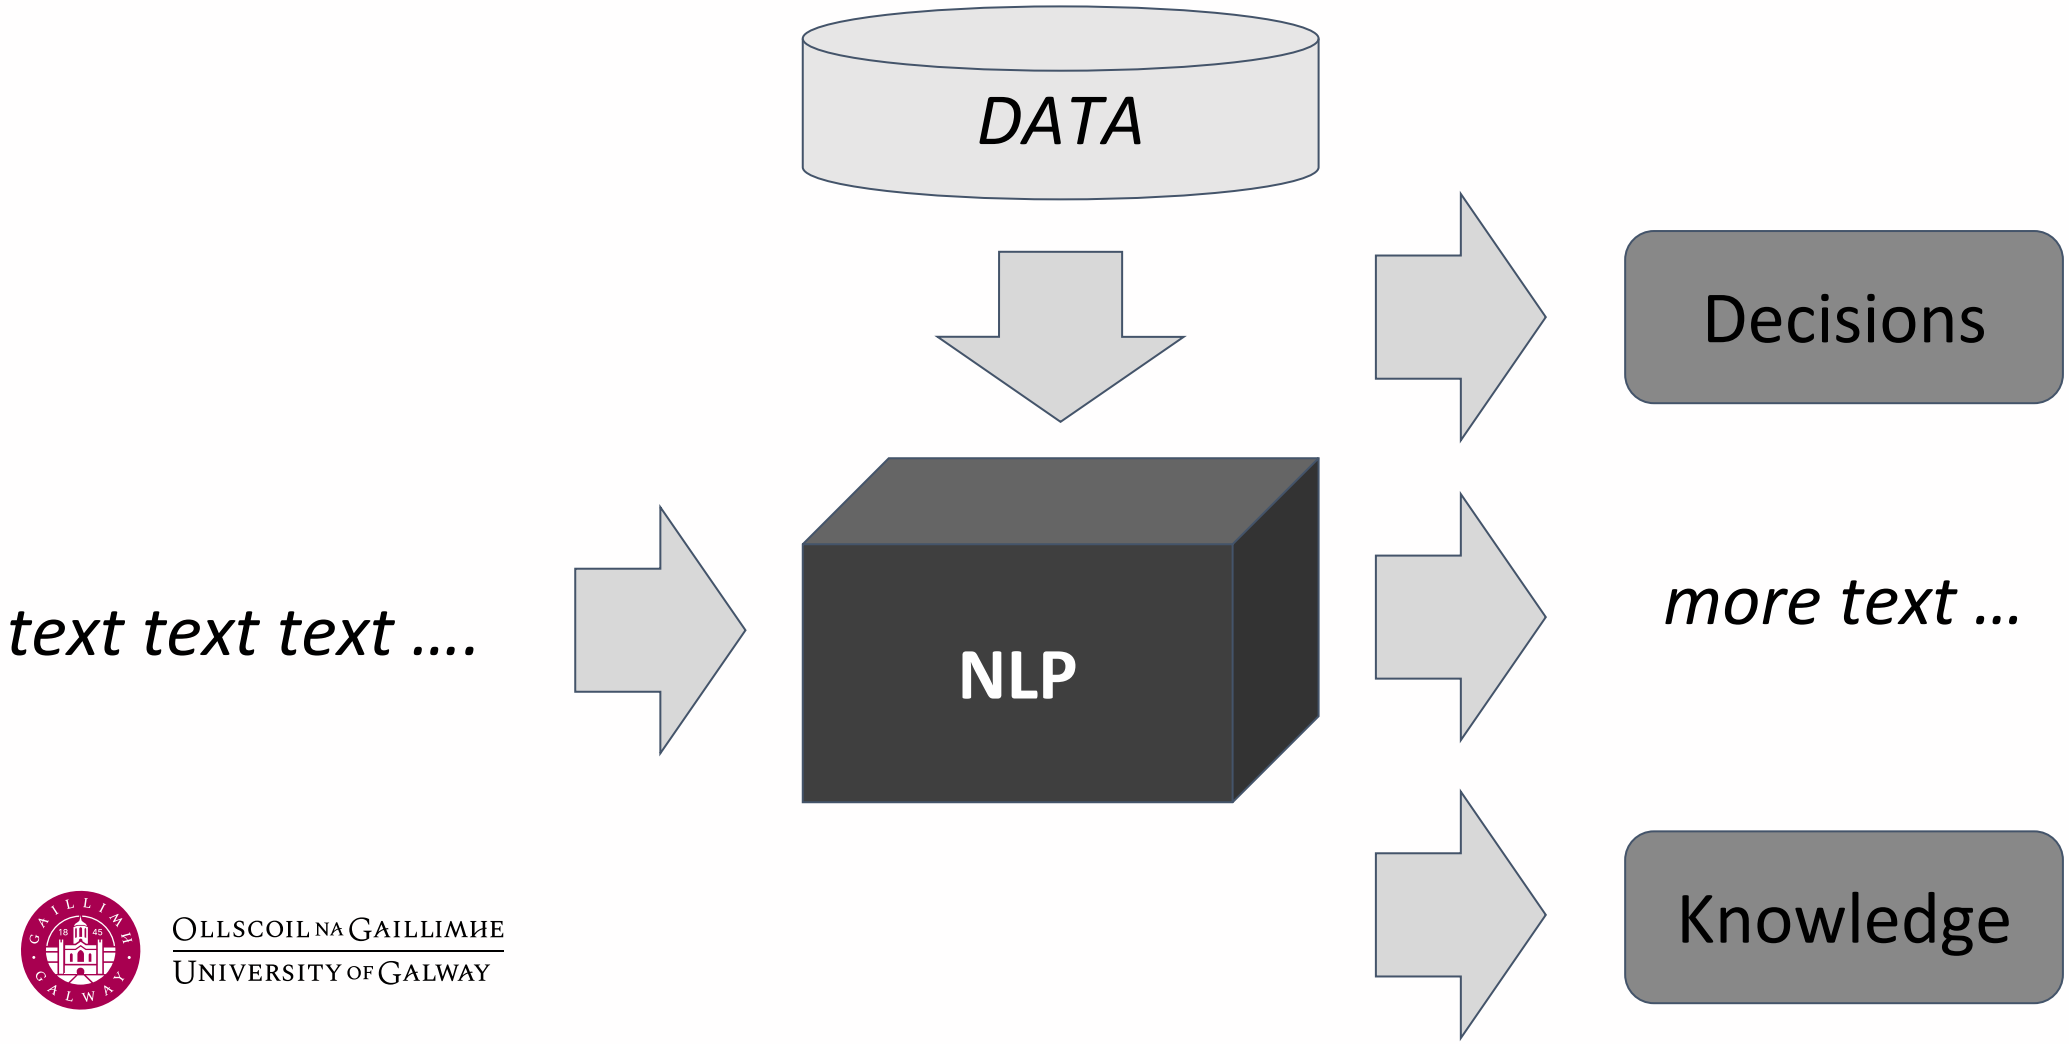

Linguistics

text text text ….

Linguistic Structure & Levels of Analysis

- **Morphology**: tokenization (MWEs), inflection, derivation, stemming, lemmatization
- **Syntax**: part-of-speech, grammar (constituency vs. dependency)
- **Semantics**: word sense, semantic roles, coreference

Language Data

- **Lexicon**: WordNet, FrameNet
- **Corpora**: annotation (data labeling), multilingual, domain-specific

Syntax: Part of Speech

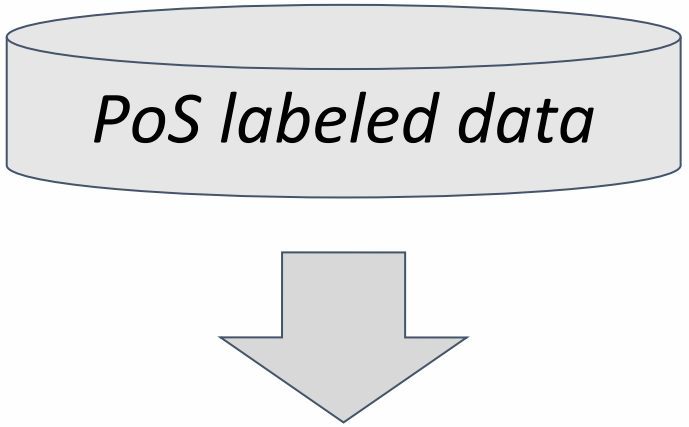

text text text ….

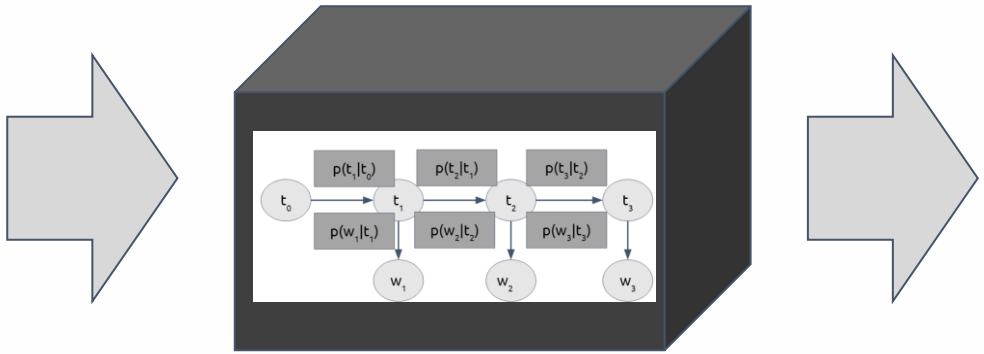

Knowledge

- Hidden Markov Model
- PoS tagged text

Syntax: Constituency / Dependency Parsing

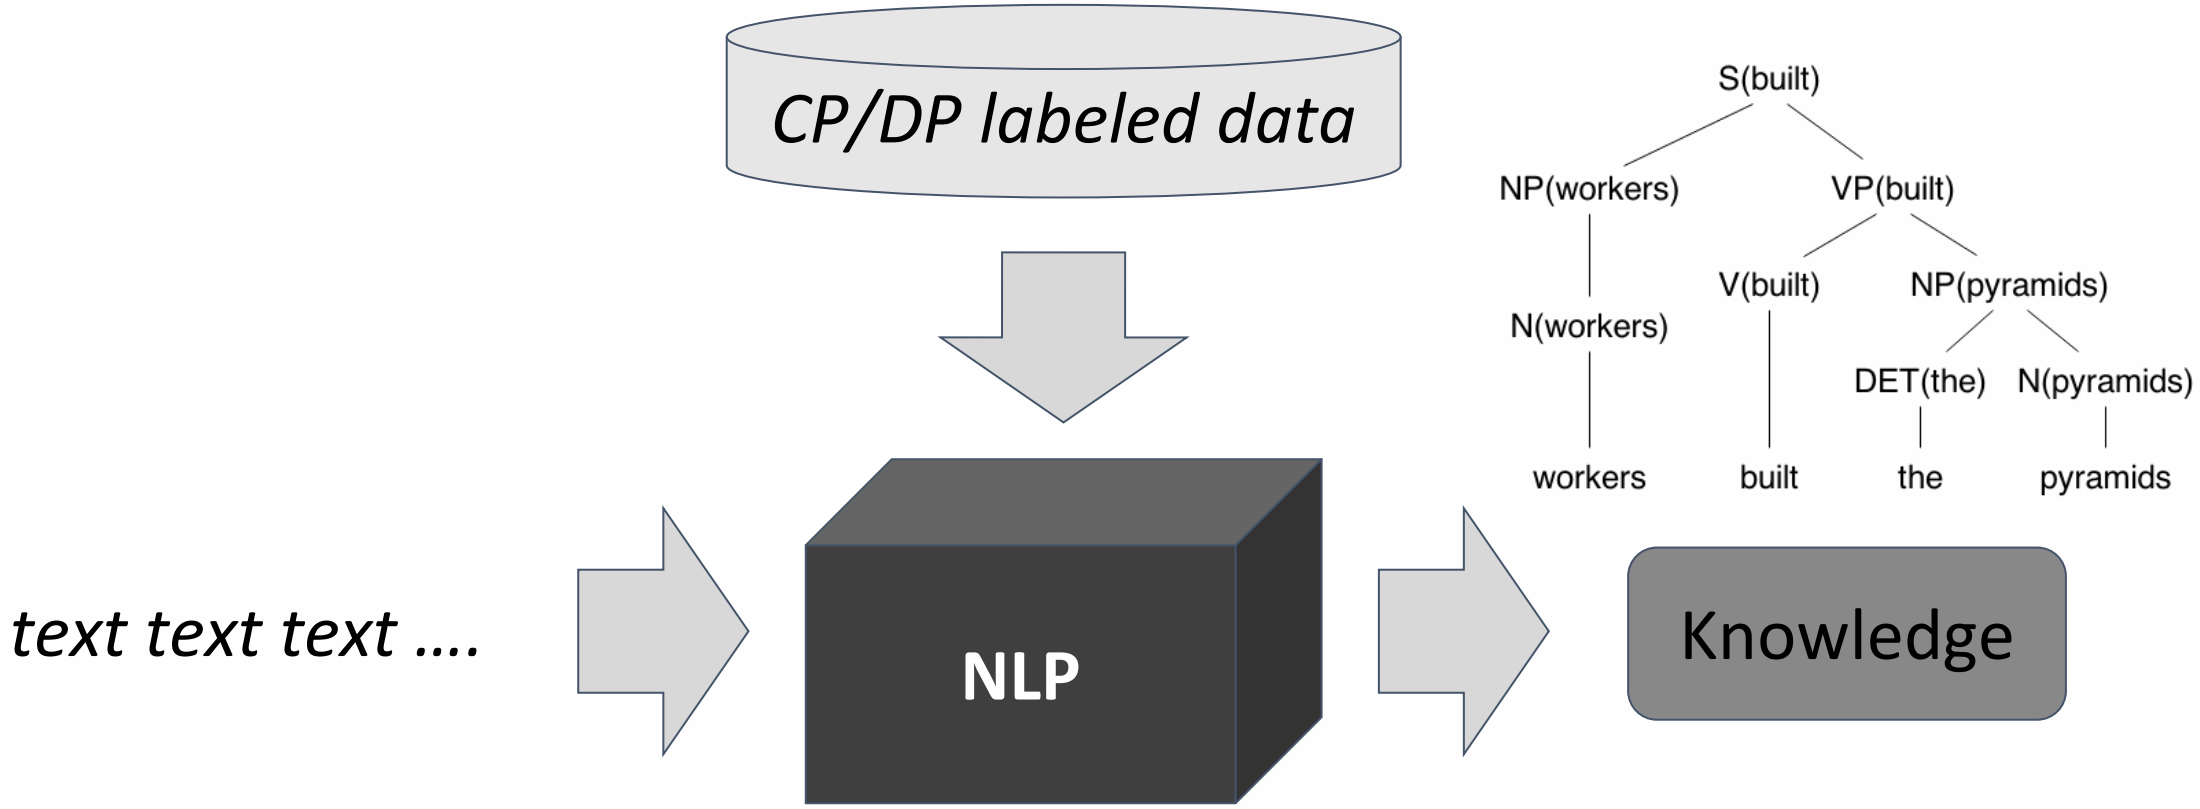

- CYK algorithm
- text annotated with parse trees

Semantics: Word Sense Disambiguation

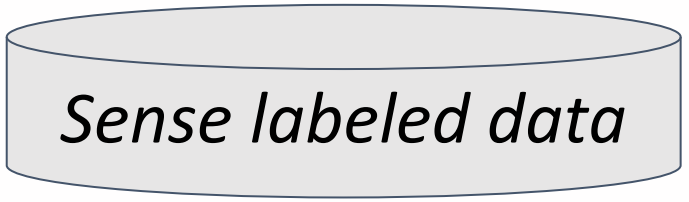

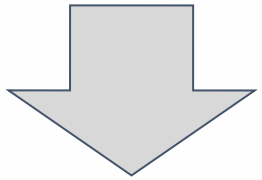

text text text ….

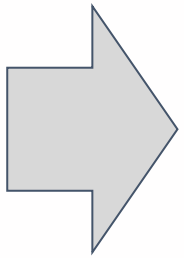

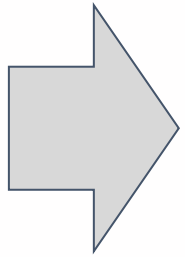

- Lesk algorithm

- word sense annotated text

Ollscoil na Gaillimhe
University of Galway

Probability

text text text ….

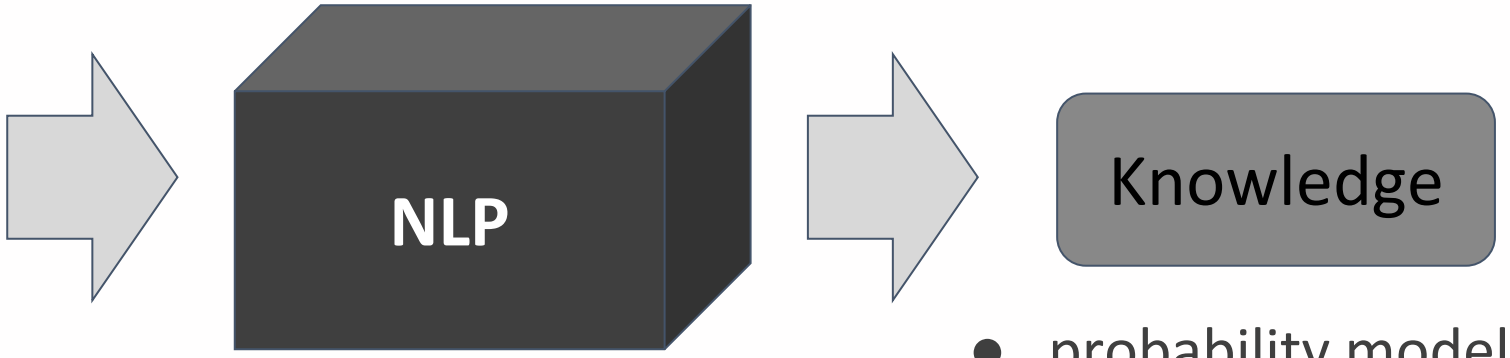

- probability model

$$P(B|A) = P(A|B)\frac{P(B)}{P(A)} \qquad TF\text{-}IDF = f_w \times \left(\log\left(\frac{N}{N_w}\right) + 1\right)$$

- estimate word probability using Naive Bayes, TF-IDF

Text Classification

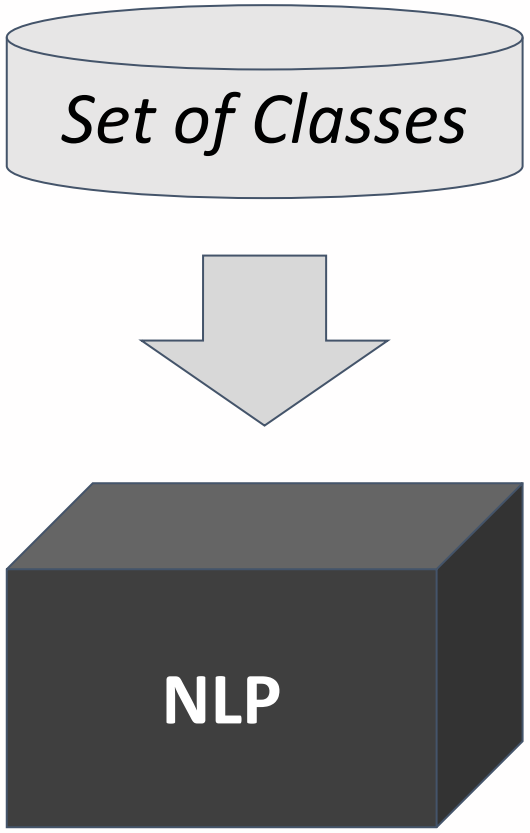

- classifier

text text text ….

- probability model

$$P(B|A) = P(A|B)\frac{P(B)}{P(A)} \qquad TF\text{-}IDF = f_w \times \left(\log\left(\frac{N}{N_w}\right) + 1\right)$$

- estimate word probability using Naive Bayes, TF-IDF

Language Modeling

text text text ….

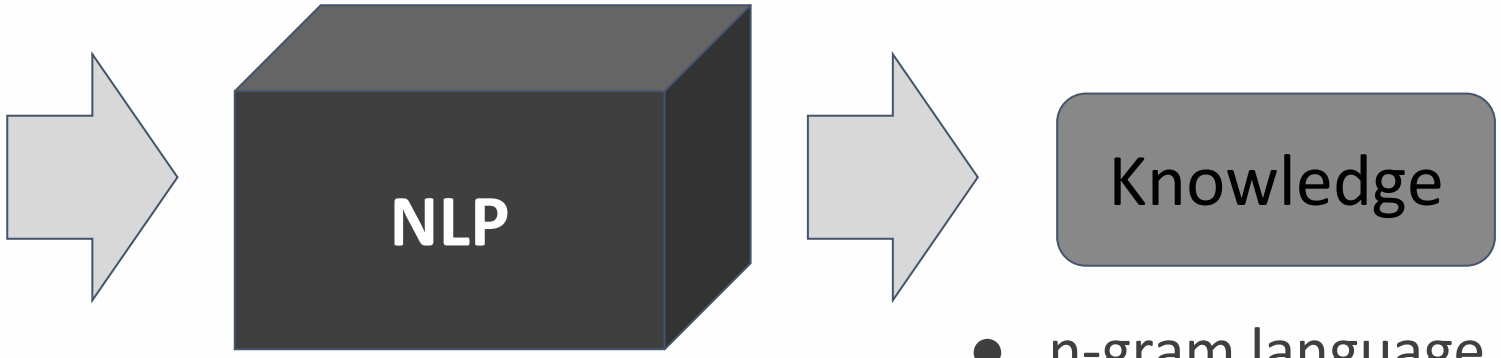

- n-gram language model

$$p(w_1 w_2 \dots w_n) = \prod_{k=1,\dots,n} p(w_k | w_{k-m+1} \dots w_{k-1})$$

- estimate the probability of a sentence using n-gram model

Vector Space

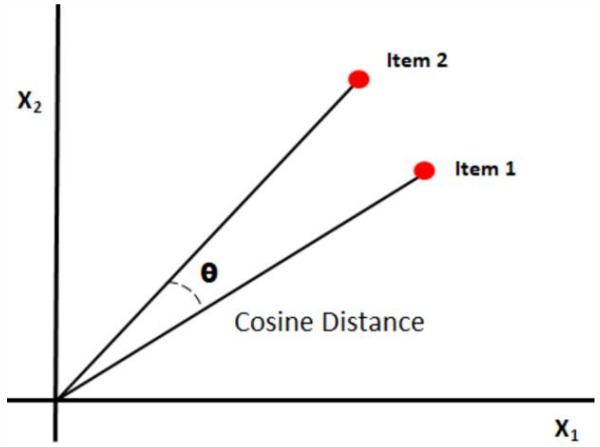

- Cosine similarity

text text text ….

- distributional model

- co-occurrence matrix

Image from https://www.oreilly.com/library/view/statistics-for-machine

Word Embeddings

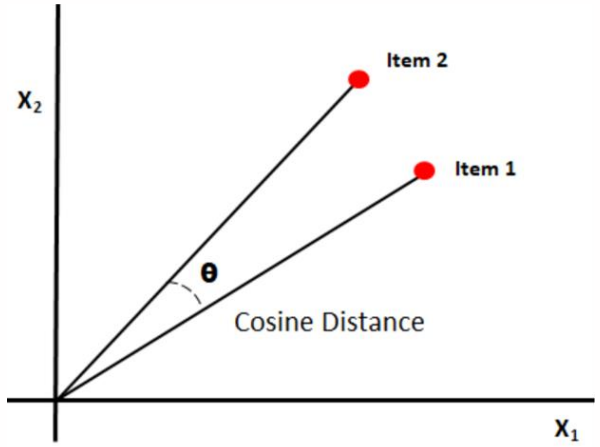

- Cosine similarity

text text text ….

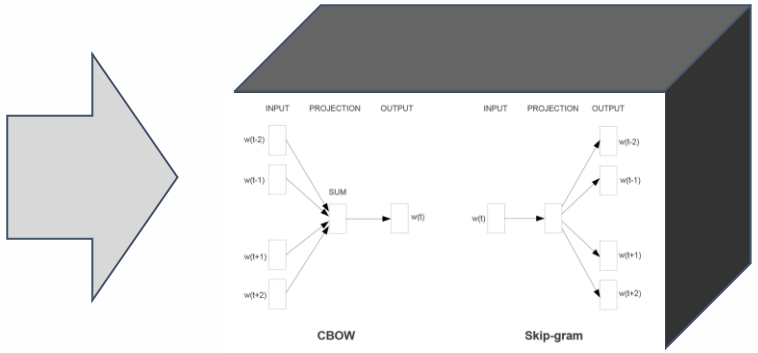

Knowledge

- word embeddings

- neural network architecture

Image from https://www.oreilly.com/library/view/statistics-for-machine

RNNs: Text Classification

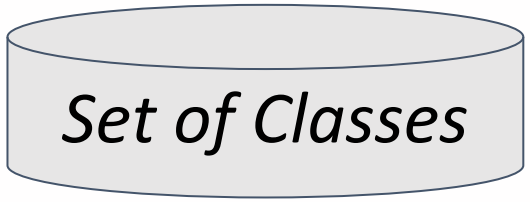

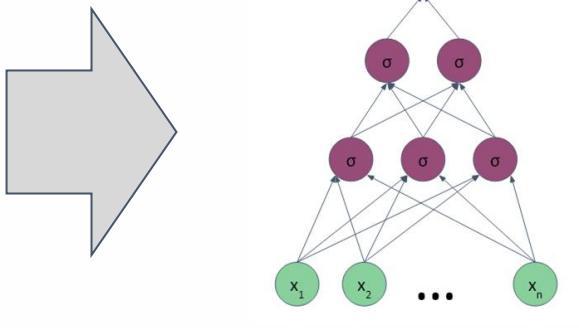

Decisions

- RNN based classifier

text text text ….

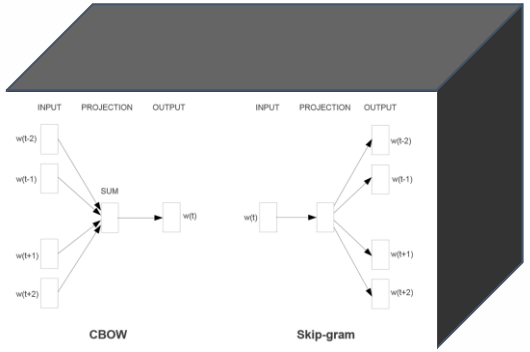

Knowledge

- word embeddings

- neural network architecture

Deep Learning: Transformers

text text text ….

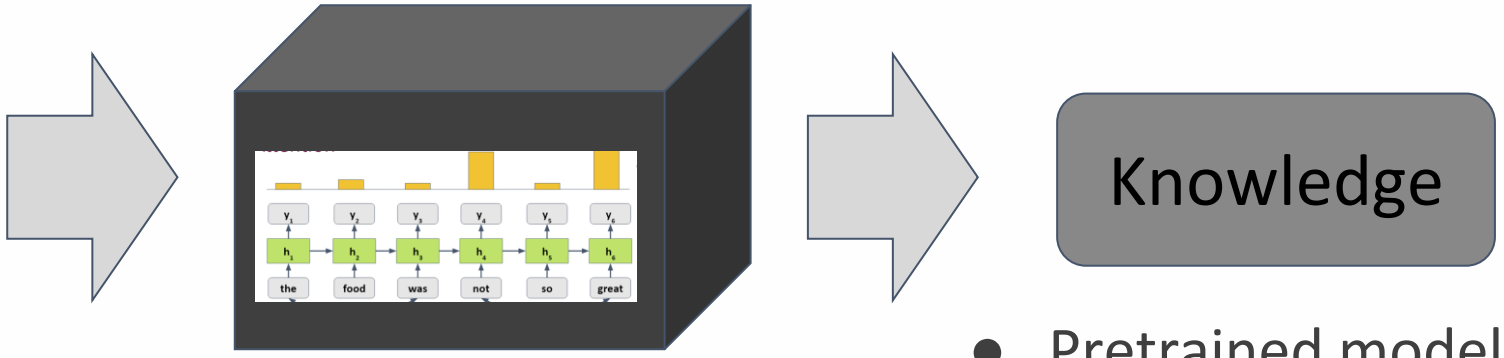

Knowledge

- Pretrained model
- Transformer with self-attention

Transformers: Text Classification

Set of Classes

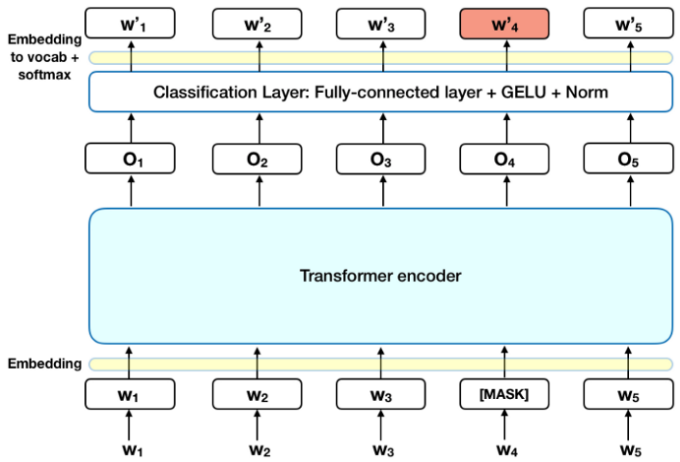

Decisions

- Classifier with fine-tuning for transfer learning

text text text ….

Knowledge

- Pretrained model

- Transformer with self-attention

Applications: Sentiment Analysis

SA labeled data

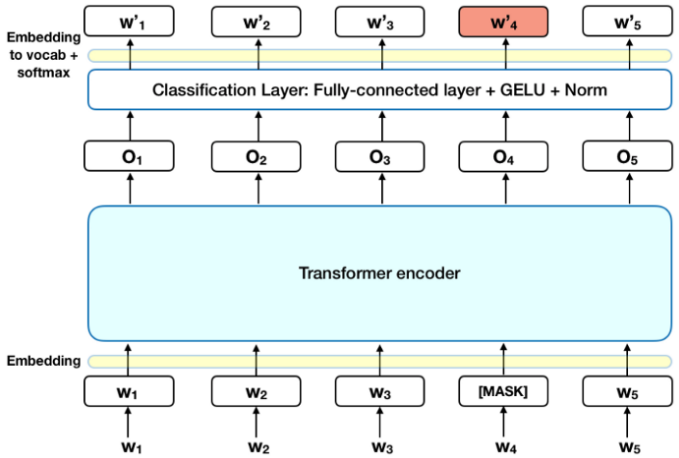

Decisions

- Sentiment analysis classifier

text text text ….

Knowledge

- Pretrained model

- Transformer with self-attention

Ethics

text text text ….

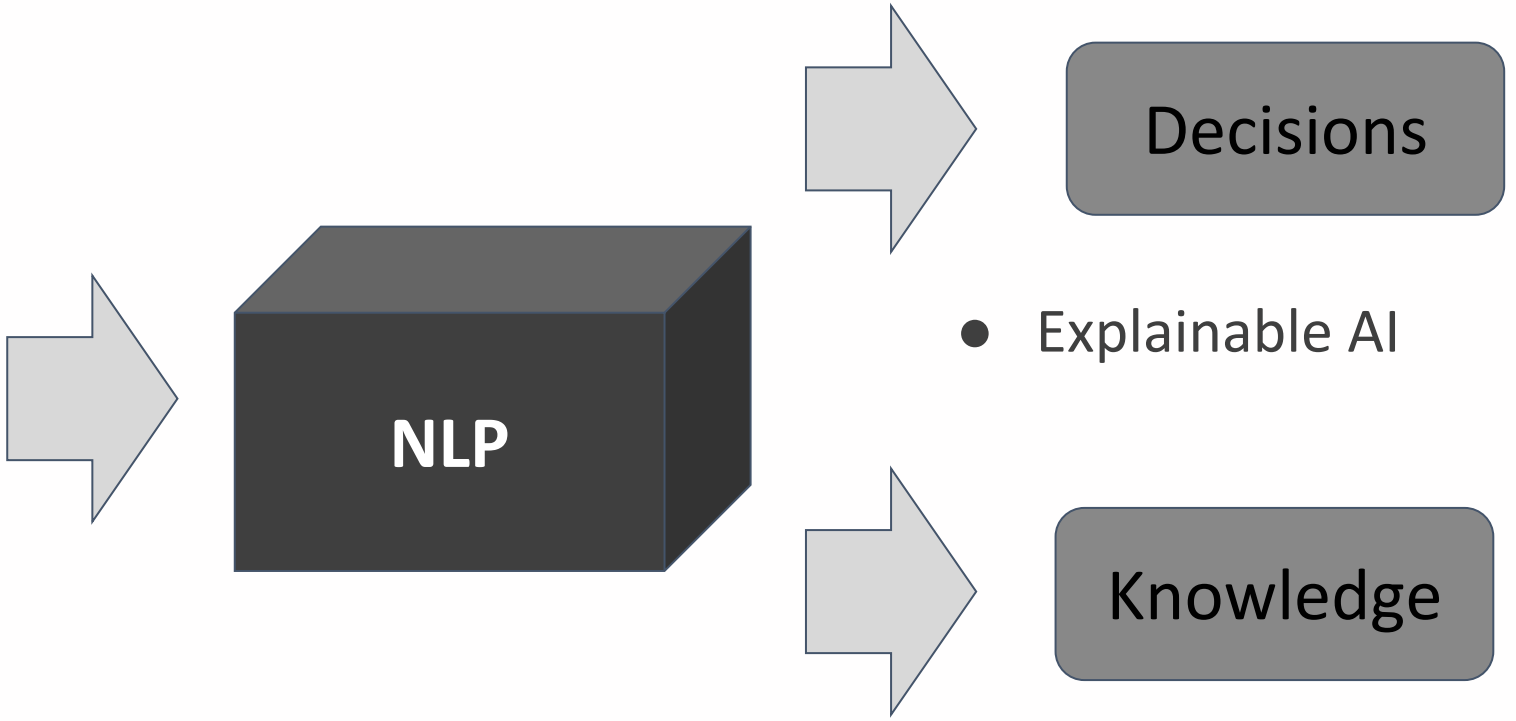

- Explainable AI

- Data privacy (GDPR)
- Data Protection Impact Assessment
- Data Statement

- Ethical, bias-aware NLP

- Trustworthy AI

Ollscoil na Gaillimhe
University of Galway

Buzz Groups

What is missing?

Ollscoil na Gaillimhe
University of Galway

NLP tasks missing
“Focus of this course”

Knowledge Extraction Tasks

Symbolic Knowledge

- Named Entity Recognition & Entity Linking
- Relation Extraction

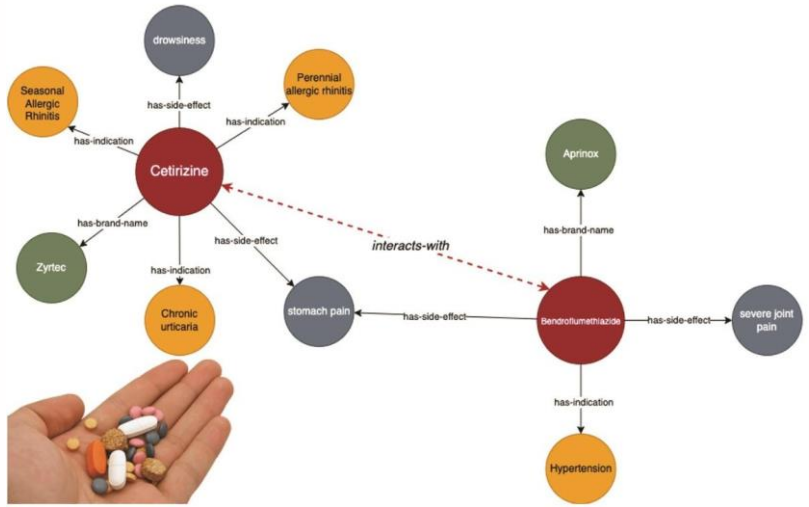

text text text ….

Knowledge

- Entities, Relations

Ollscoil na Gaillimhe
University of Galway

Image source: https://timesofmalta.com/articles/view/identifying-drug-interactions-using-knowledge-graphs.936927

Knowledge Extraction - Definition

Creation of knowledge from **unstructured (textual, language) data**

Extracted knowledge must **facilitate inferencing**

Requires **reuse and/or generation of formal knowledge**

Ollscoil na Gaillimhe
University of Galway

Adapted from *https://en.wikipedia.org/wiki/Knowledge_extraction*

Knowledge Extraction - Example

"Naomi Carey, who is the director of Hutchinson Care Homes said she is currently only able to operate at 85% capacity."

"Galway GP Martin Daley, former president of the Irish Medical Organisation, gave a statement today to this effect."

Entities

PERSON:	*Martin Daley, Naomi Carey*
ORGANISATION:	*Irish Medical Organisation, Hutchinson Care Homes*
CITY:	*Galway*

Relations

Martin Daley →	at-organisation	← *Irish Medical Organisation*
Naomi Carey →	at-organisation	← *Hutchinson Care Homes*
Martin Daley →	has-occupation	← *GP*
Naomi Carey →	has-occupation	← *former president*

Ollscoil na Gaillimhe
University of Galway

Opinion Mining Tasks

Emotion Analysis

Aspect Mining

Figurative Language Processing

Emotion Analysis

Emotion Analysis

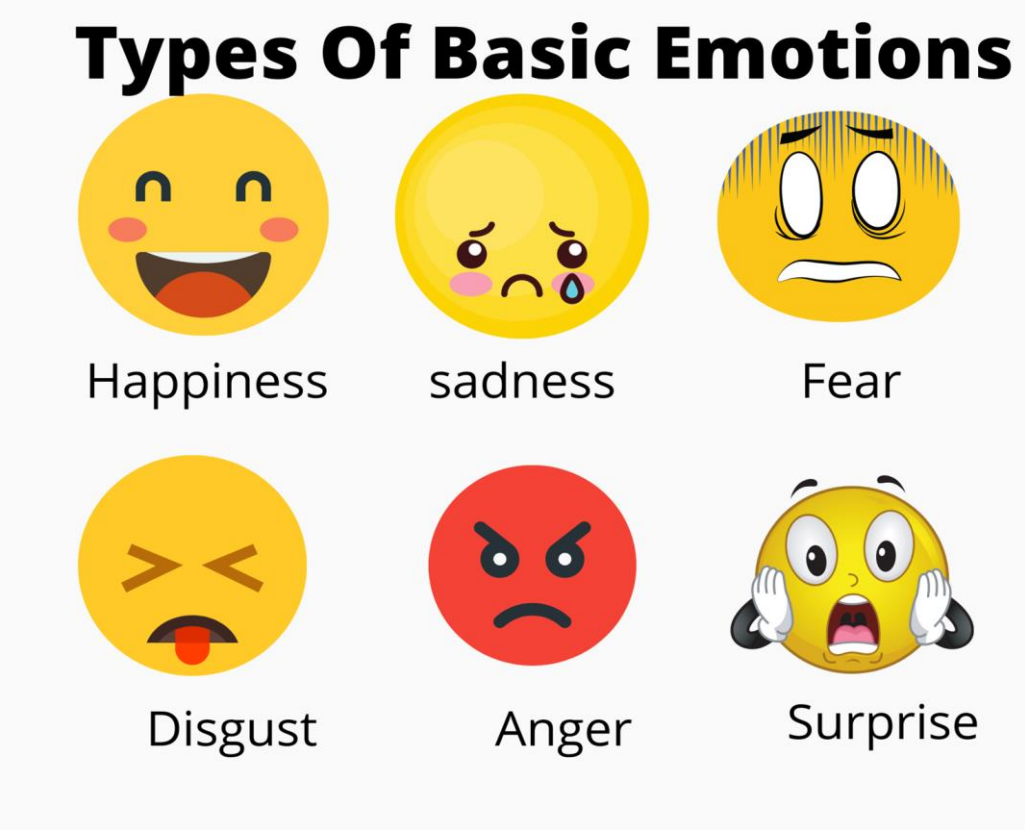

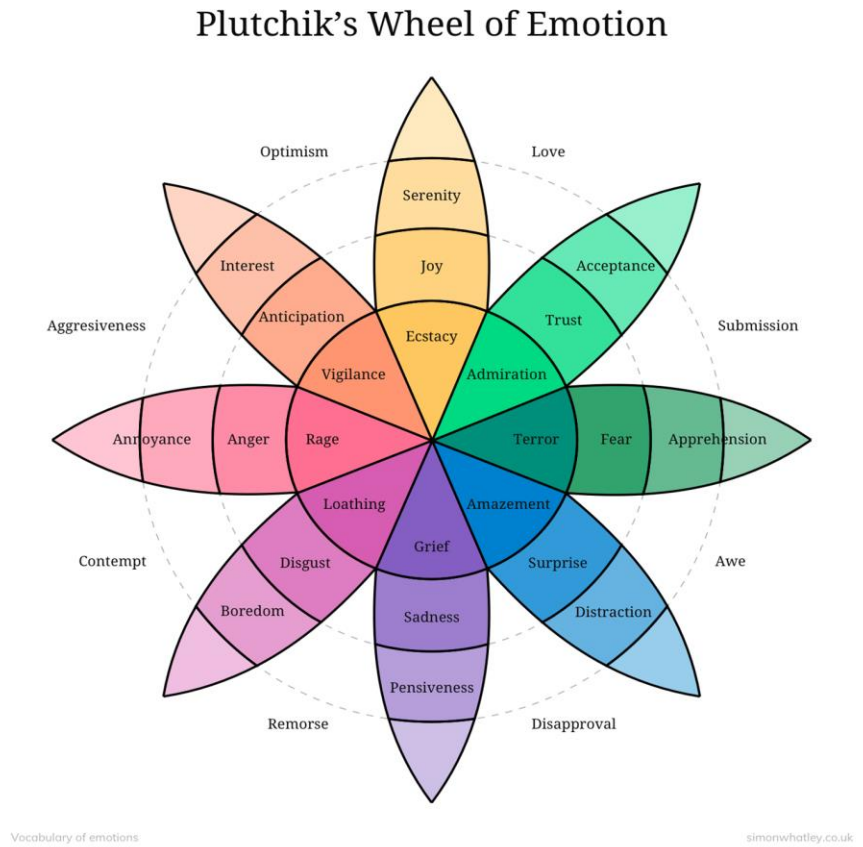

Emotion Analysis

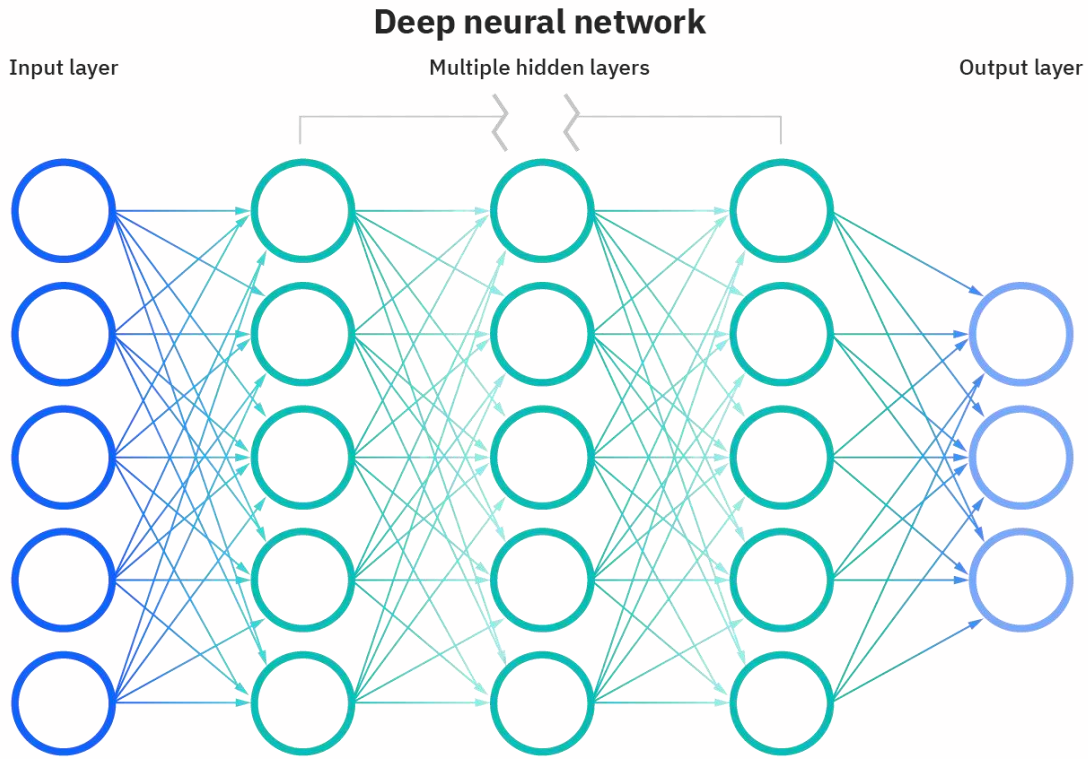

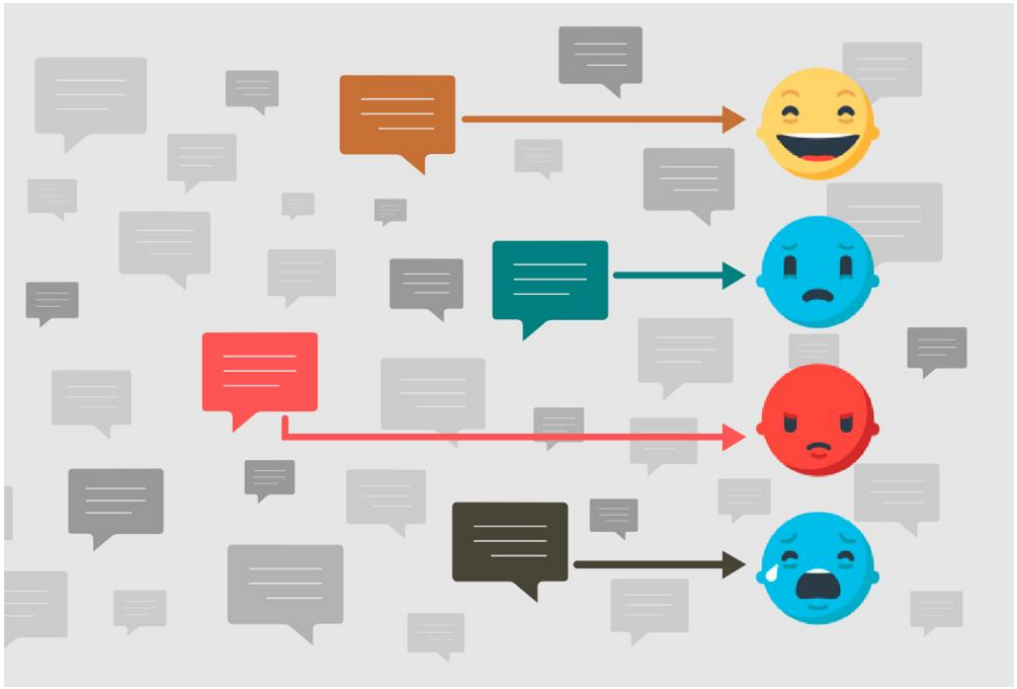

Aspect Mining

*"The camera's focus was **bad**, but has a **great** size and is **easy-to-use**."*

Aspect Mining

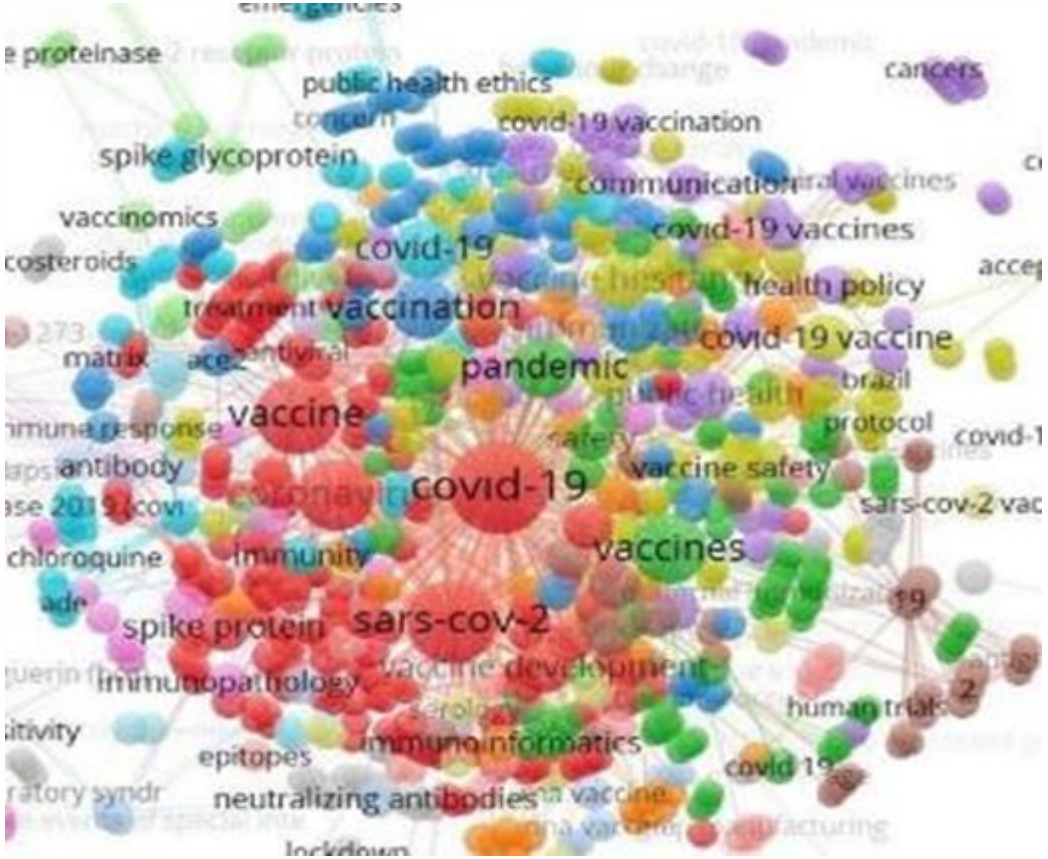

Figurative Language Processing

Idiom (noun): a group of words established by usage as having a meaning not deducible from those of the individual words

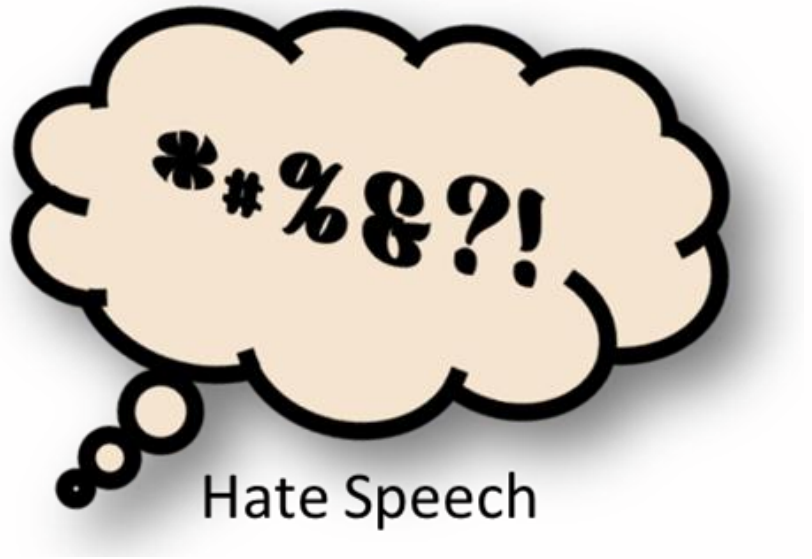

Figurative Language Processing

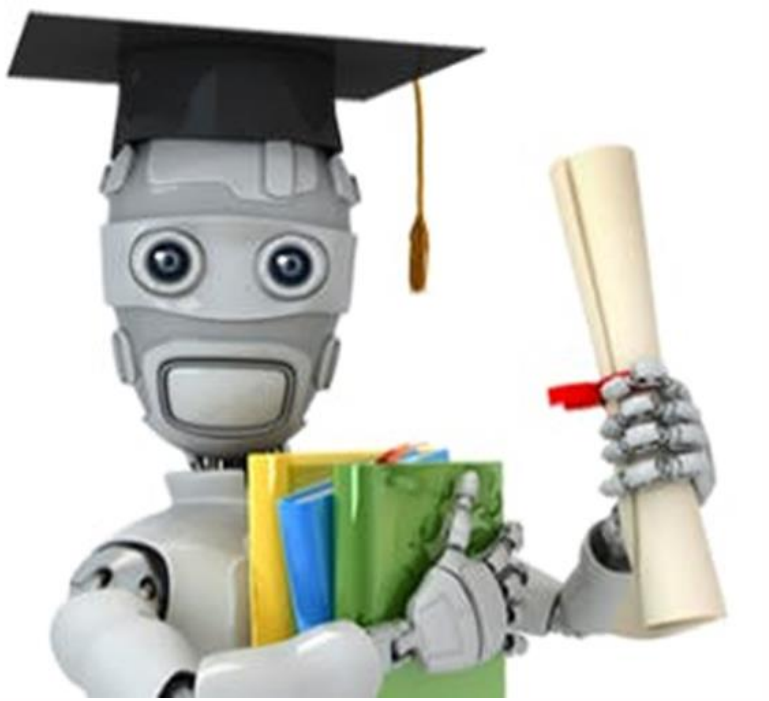

Language Generation Tasks

Machine Translation

Data-to-Text Generation

Dialog System Development

Ollscoil na Gaillimhe
University of Galway

Language Generation: Machine Translation

English text ….

Nederlandse tekst …

Machine Translation

gerade zu diesem Stamm gehören ||| belong just to these families ||| 0.390442 2.9025
gerade zu diesem Stamm gehören ||| belong just to these ||| 0.390442 2.9025e-15 0.2
gerade zu diesem Stamm gehören ||| them belong just to these ||| 0.390442 2.9025e-
gerade zu diesem Stamm ||| belong just to these families ||| 0.390442 1.5563e-11 0.2
gerade zu diesem Stamm ||| belong just to these ||| 0.390442 1.5563e-11 0.260295 9.
gerade zu diesem Stamm ||| them belong just to these ||| 0.390442 1.5563e-11 0.260
gerade zu diesem ||| belong just to ||| 0.390442 1.33531e-08 0.260295 4.9937e-09 2.
gerade zu diesem ||| of them belong just to ||| 0.390442 1.33531e-08 0.260295 4.79
gerade zu diesem ||| them belong just to ||| 0.390442 1.33531e-08 0.260295 4.6096
gerade zu ||| belong just to ||| 0.390442 1.8515e-05 0.260295 4.9937e-09 2.718 |||
gerade zu ||| of them belong just to ||| 0.390442 1.8515e-05 0.260295 4.79751e-13
gerade zu ||| them belong just to ||| 0.390442 1.8515e-05 0.260295 4.60968e-12 2.7

Statistical Machine Translation

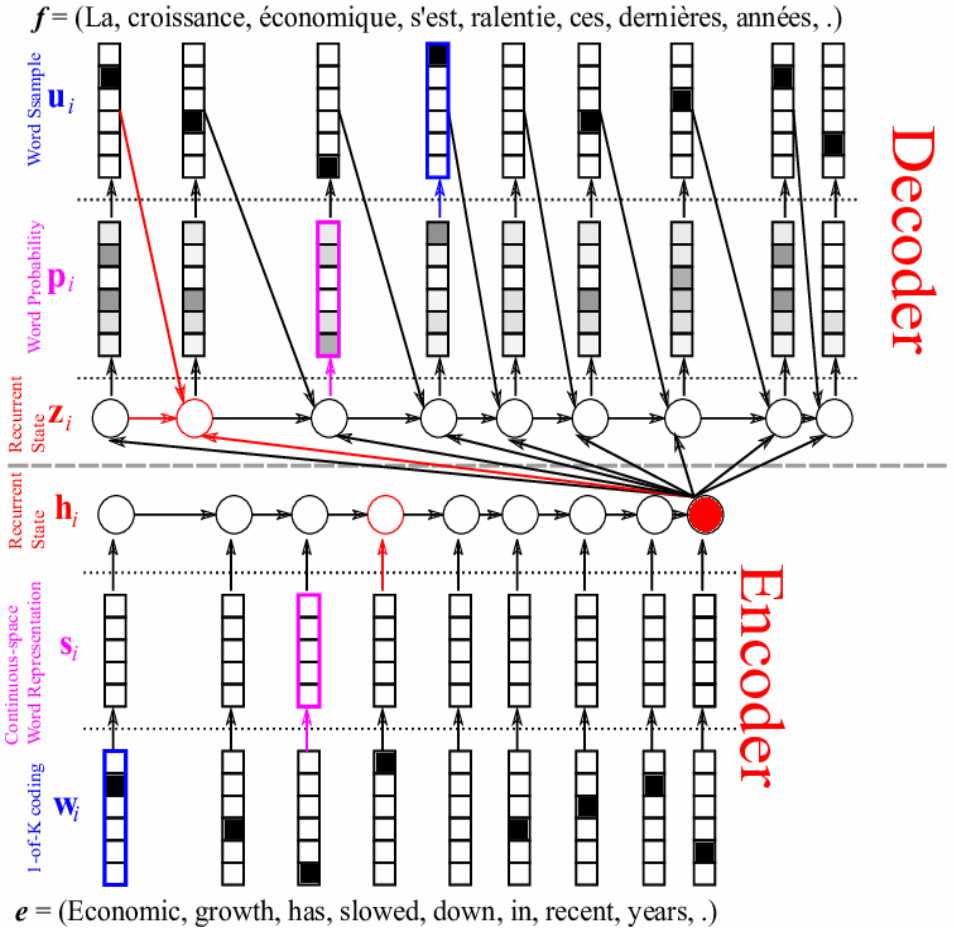

Neural Machine Translation

Language Generation: Data-to-Text

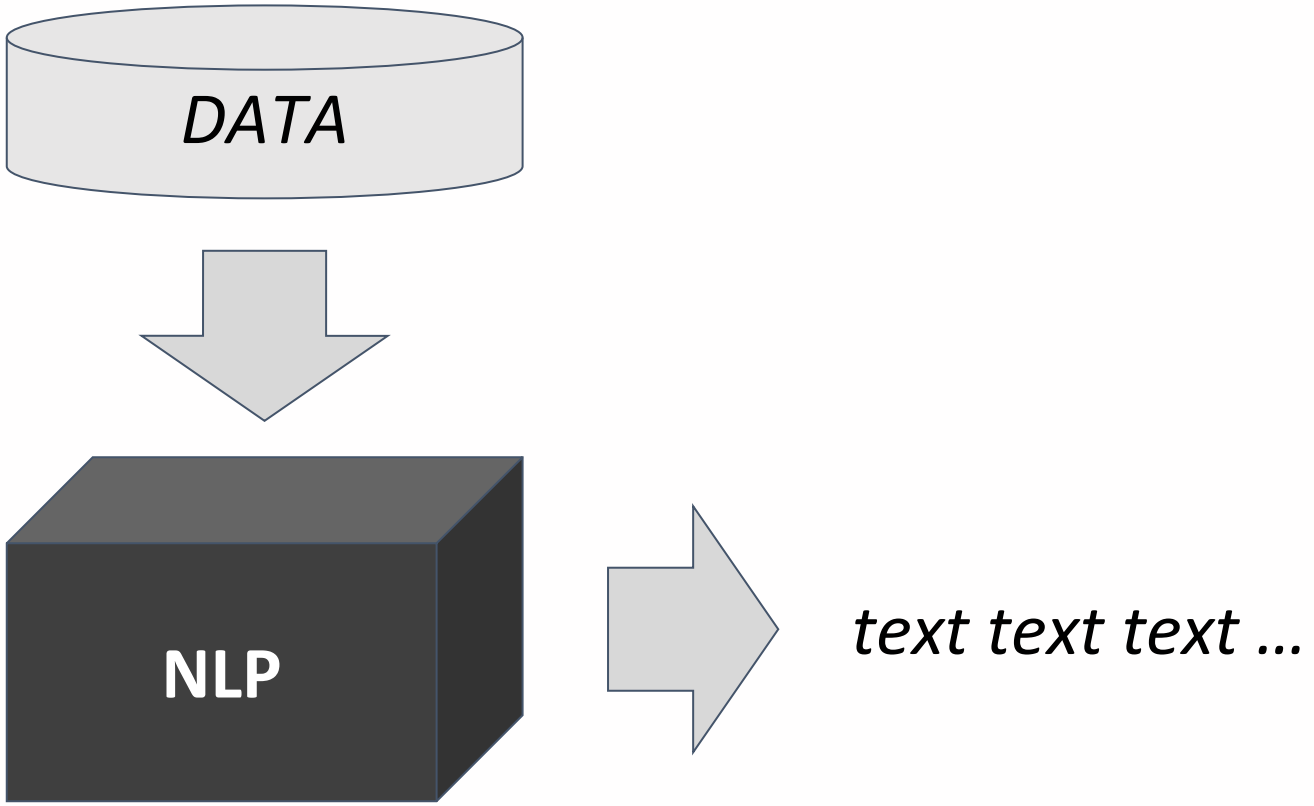

Data-to-Text Generation

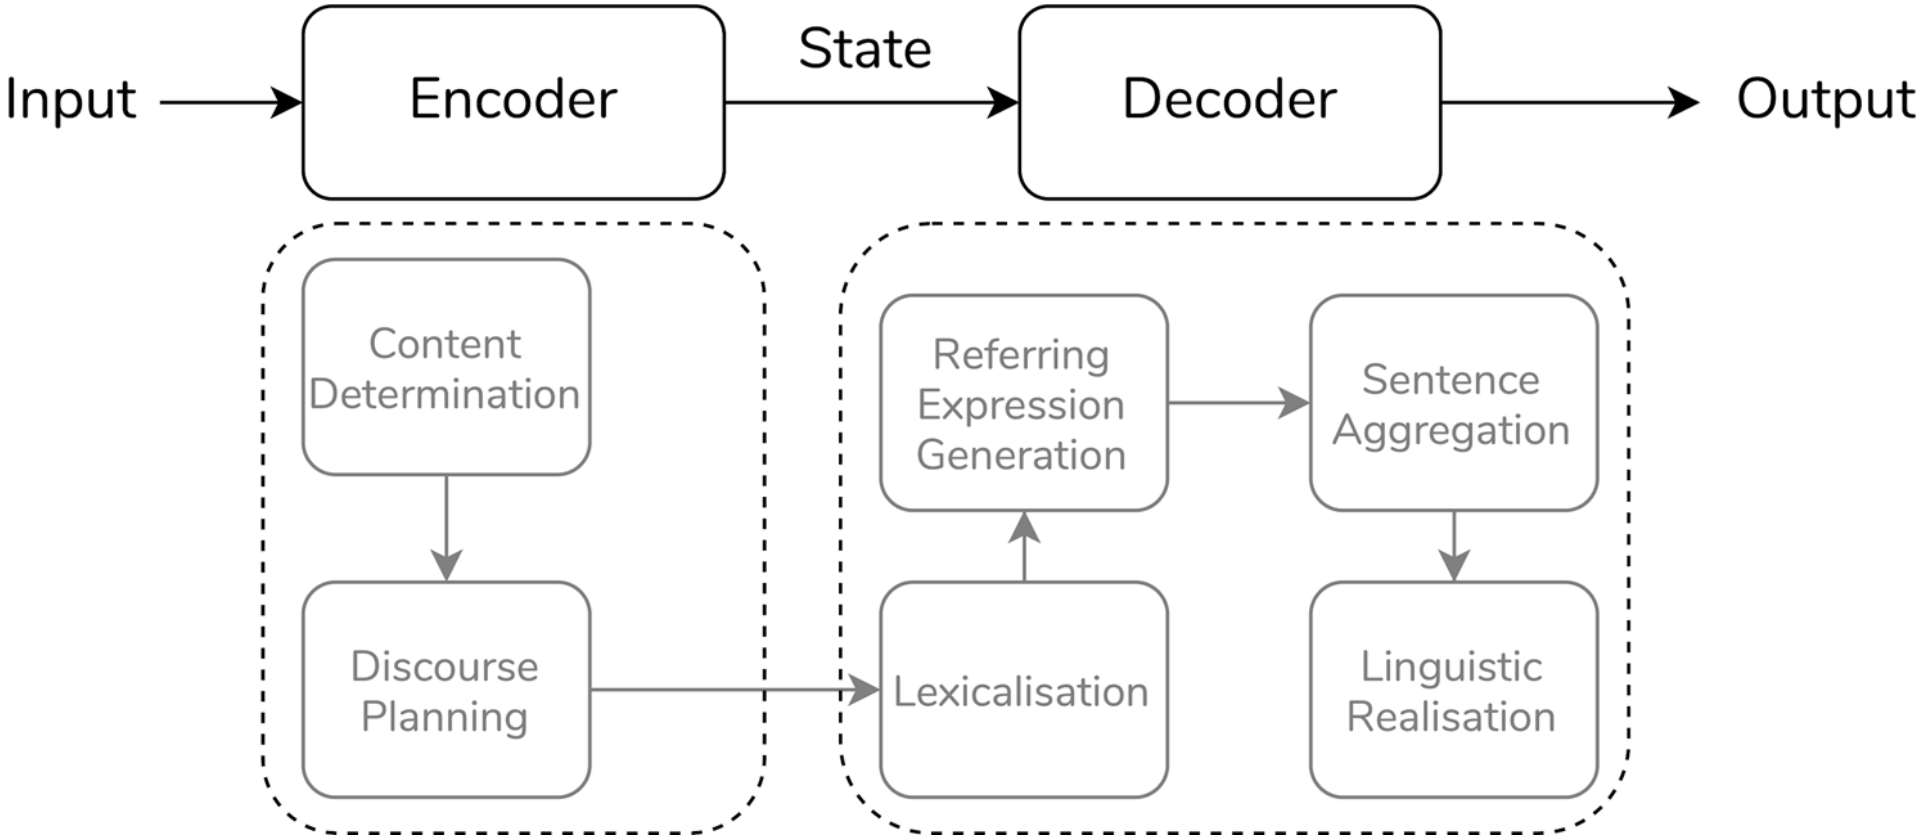

Image by Nivranshu Pasricha

Language Generation: Dialog Systems

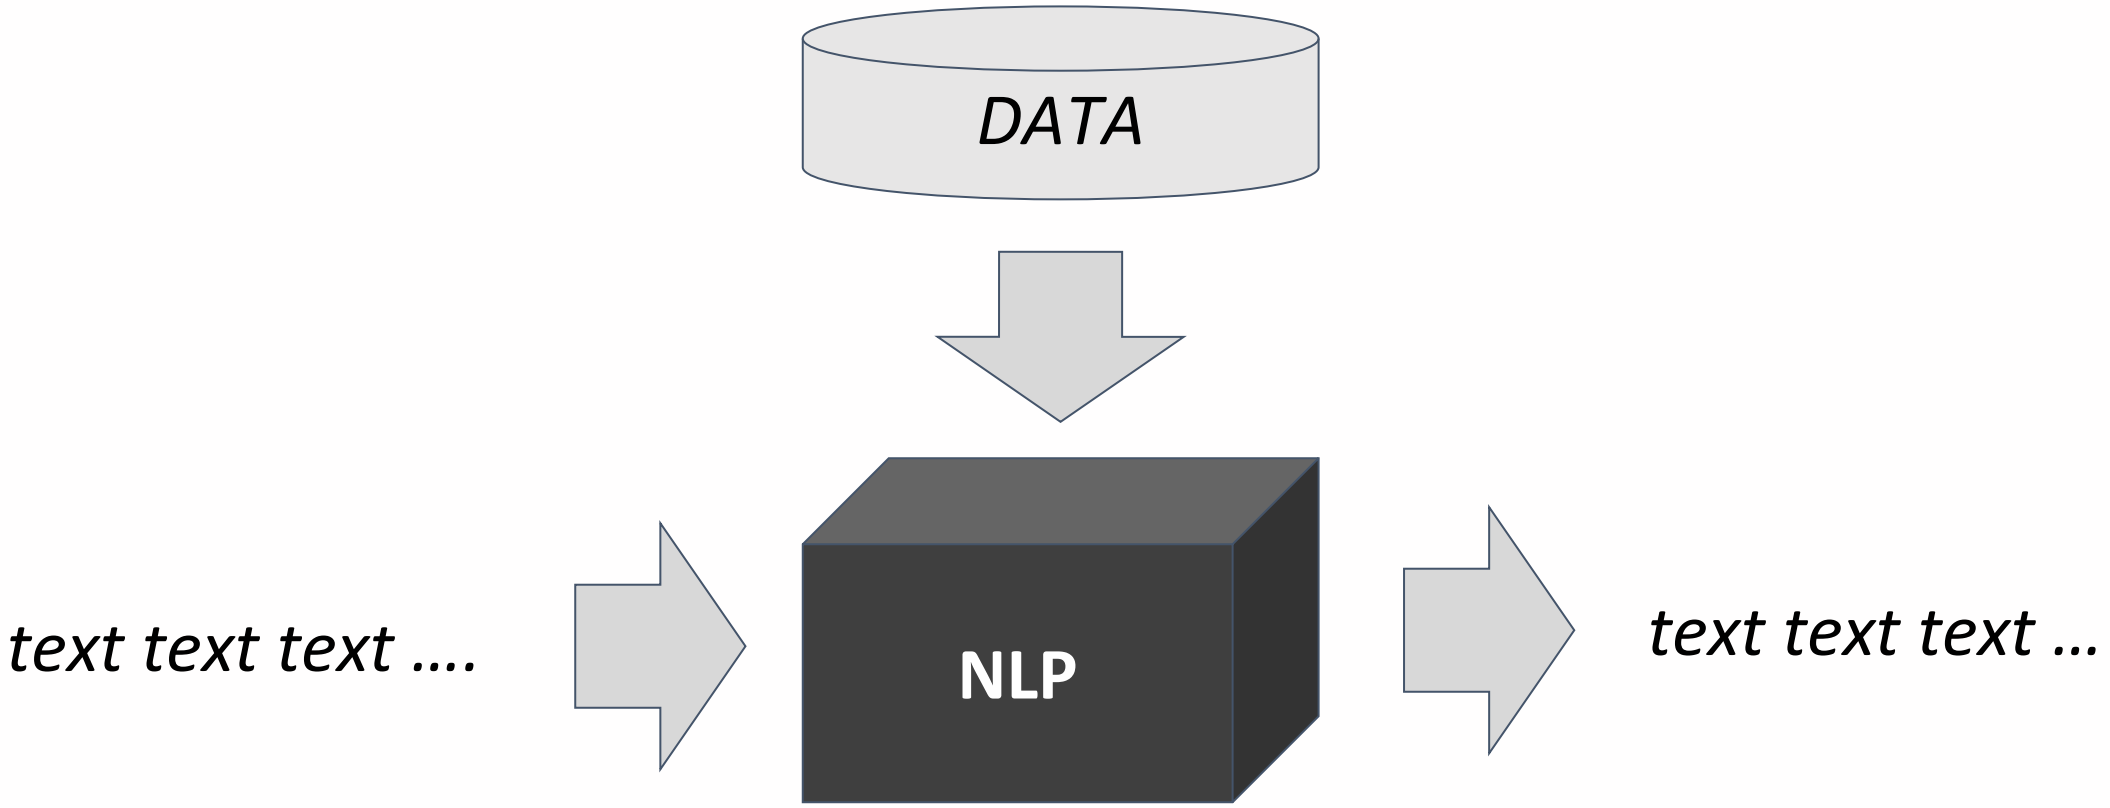

Dialog System Development

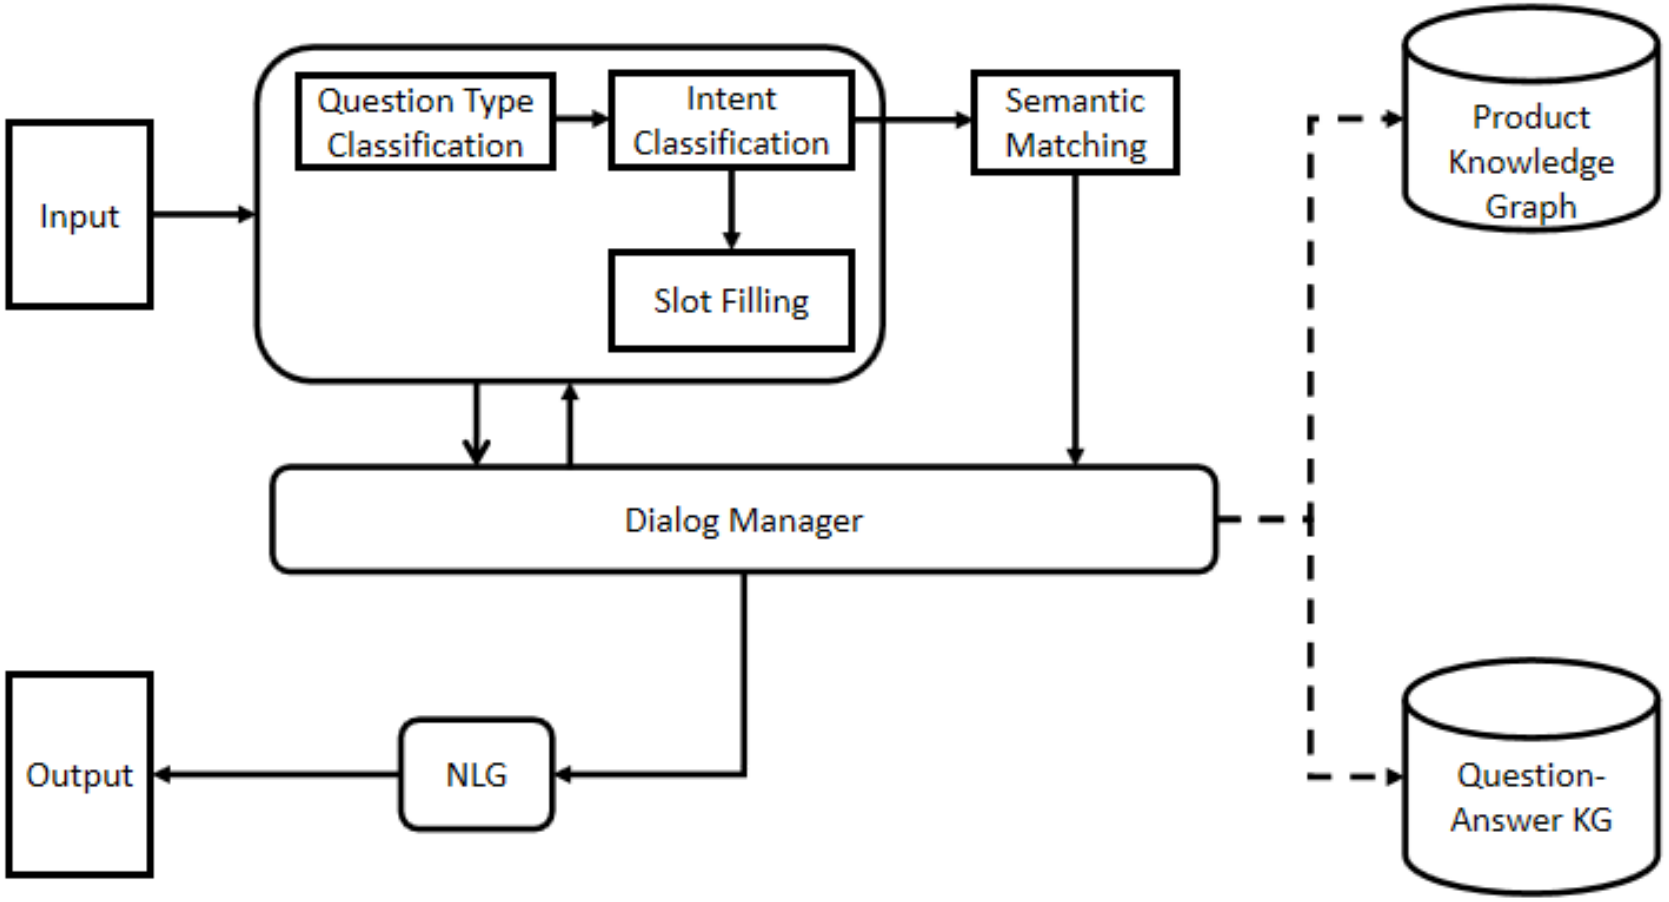

Ollscoil na Gaillimhe
University of Galway

Image from: Zhao, G., Zhao, J., Li, Y., Alt, C., Schwarzenberg, R., Hennig, L., Schaffer, S., Schmeier, S., Hu, C. and Xu, F., 2019. MOLI: Smart Conversation Agent for Mobile Customer Service. *Information*, 10(2), p.63.

Ollscoil na Gaillimhe
University of Galway

NLP in industry

University
*of*Galway.ie

How Google uses NLP to better understand search queries, content

Learn the role that natural language processing plays in making Google search even more semantic and context-based.

Olaf Kopp on August 23, 2022 at 6:00 am | Reading time: 10 minutes

Natural language processing opened the door for semantic search on Google.

SEOs need to understand the switch to entity-based search because this is the future of Google search.

Ollscoil na Gaillimhe
University of Galway

https://searchengineland.com/how-google-uses-nlp-to-better-understand-search-queries-content-387340

Google Index & Knowledge Graph

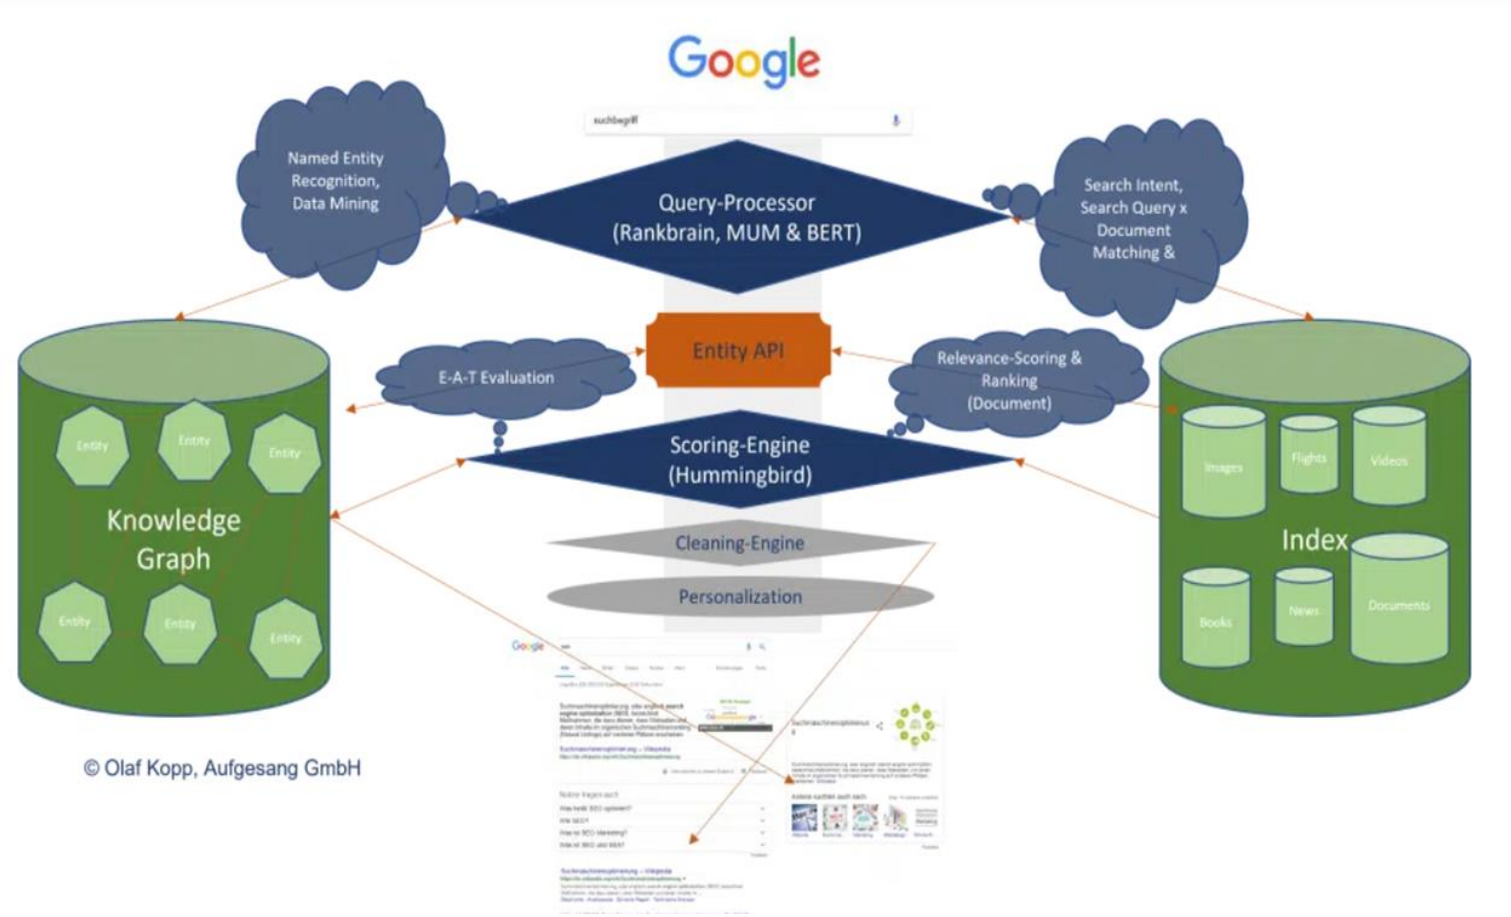

https://searchengineland.com/how-google-uses-nlp-to-better-understand-search-queries-content-387340

“Use of NLP in Google Search”

According to Olaf Kopp of Aufgesang GmbH (article in Search Engine Land), Google Search uses NLP for the following:

- Interpretation of search queries.
- Classification of subject and purpose of documents.
- Entity analysis in documents, search queries and social media posts.
- Generating featured snippets and answers in voice search.
- Interpretation of video and audio content.
- Expansion and improvement of the Knowledge Graph.

https://searchengineland.com/how-google-uses-nlp-to-better-understand-search-queries-content-387340

Ollscoil na Gaillimhe
University of Galway

Buzz Groups

Which NLP tasks are needed for each of these steps?

Interpretation of search queries

“calories in bar of chocolat”

- Relation Extraction, Parsing, Part-of-Speech, WSD

“Biden”

- Entity Linking

“great movies” “scary movie”

- Sentiment / Emotion Analysis

Classification of subject, purpose of documents

legal, health, … documents

- Text Classification, Probability, Language Modeling, Word Embeddings

documents about concept XYZ

- Concept Extraction

Entity analysis in documents, queries and posts

identify names of people, locations, things

- NER, Entity Linking, Concept/Taxonomy Extraction

identify relations between people, locations, things

- Relation Extraction, Parsing, Semantic Role Labeling

resolution of pronouns

- Coreference Resolution

Generating featured snippets and answers

summarize one or more retrieved documents

- Text Summarization

generate a specific answer

- Natural Language Generation, Dialog

translate a text

- Machine Translation

Ollscoil na Gaillimhe
University of Galway

Interpretation of video and audio content

retrieve relevant videos for a search query

- Multimodal Analysis

speech interaction

- Speech-to-Text, Dialog

Expansion/improvement of Knowledge Graph

include new entities

- Entity Linking, Concept Extraction, Taxonomy Extraction/Extension

include new or update existing relations

- Relation Extraction, Knowledge Graph Completion

Lab of this week

Intro to base NLP tools and methods

Tools:

- pandas
- Hugging Face Datasets
- scikit-learn
- Hugging Face Transformers

Dataset:

- IMDB reviews: https://huggingface.co/datasets/imdb