

CT5132/CT5148 Week 11 Exercises

James McDermott

University of Galway

Week 11 Exercises

The exercises and solutions are extracted from the lecture slides.

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse
## v ggplot2 3.2.1      v purrr  0.3.2
## v tibble  2.1.3      v dplyr  0.8.3
## v tidyr   1.0.0      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0
## -- Conflicts ----- tidyverse_confli
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

Exercises (R Basics)

- 1 Write the Factorial function in R, eg `fact(5)` gives 120.
- 2 Given `x <- "John"`, calculate the length in characters of `x`. Use `nchar()`.
- 3 Given `xs <- c("John", "Paul", "George", "Ringo")`, calculate the length of each name, using vectorisation (not a for-loop).
- 4 Calculate whether each name is shorter than 5 characters.
- 5 Index `xs` to keep just the names shorter than 5 characters.
- 6 Write a function which unit-norms a vector, ie normalises it so that the vector length equals 1. Eg `unit_norm(c(10, 10, 10, 10))` gives 0.5 0.5 0.5 0.5.
- 7 Write a function which standardises a vector, ie gets the z-score, ie maps it to have mean 0 and standard deviation 1. Eg `z_score(c(10, 6, 12, 12))` gives 0.0000000 -1.4142136 0.7071068 0.7071068.

Solutions (R Basics)

```
fact <- function(n) { # Exercise 1
  if (n <= 1) {
    1 # remember, no return statement!
  } else {
    n * fact(n-1)
  }
}
fact(5)

## [1] 120
```

```
x <- "John"
nchar(x) # Exercise 2
```

```
## [1] 4
```

```
xs <- c("John", "Paul", "George", "Ringo")
nchar(xs) # Exercise 3
```

```
## [1] 4 4 6 5
```

```
nchar(xs) < 5 # Exercise 4
```

```
## [1] TRUE TRUE FALSE FALSE
```

```
xs[nchar(xs) < 5] # Exercise 5
```

```
## [1] "John" "Paul"
```

```
unit_norm <- function(x) { # Exercise 6
  x / sqrt(sum(x**2))
}
unit_norm(c(10, 10, 10, 10))
```

```
## [1] 0.5 0.5 0.5 0.5
```

```
z_score <- function(x) { # Exercise 7
  (x - mean(x)) / sd(x)
}
z_score(c(10, 6, 12, 12))
```

```
## [1] 0.0000000 -1.4142136 0.7071068 0.7071068
```

Exercises (Tidy Data)

- 1 Recall our experiment on running time for sorting an array of different sizes. The original data (before we added extra columns) is available in `data/sort_times_original.csv`. Read it in to a tibble. (You might need to set the working directory first.)
- 2 Use `glimpse` to take a look. What types do the columns have?
- 3 In what way is this *not* tidy data? Use `gather` to fix it. Hint: the result should have shape 50×3 with columns `n`, `run_number`, `run_time`.
- 4 It would be nicer if `run_number` was just an integer, eg 0, instead of `run0`. Use `separate` to split it into two parts. Hint: use `into=c("dummy", "run_number")`.
- 5 Look again at the result. We don't need that “dummy” column. Use `NA` to omit it. Hint: see `?separate` for help on `into`.
- 6 Look again – `run_number` is still not an integer! Fix this. Hint: `separate` can guess the correct type to convert to, but see `?separate` again to see how to ask it to.
- 7 Write it to a file `data/sort_times_tidy.csv` using `write_csv()`.

Solutions (Tidy Data)

Exercise 1

```
d <- read_csv("data/sort_times_original.csv")

## Parsed with column specification:
## cols(
##   n = col_double(),
##   run0 = col_double(),
##   run1 = col_double(),
##   run2 = col_double(),
##   run3 = col_double(),
##   run4 = col_double()
## )
```

Exercise 2

```
glimpse(d) # All columns of type `dbl`, which is ok
```

```
## Observations: 10
```

```
## Variables: 6
```

```
## $ n      <dbl> 1e+06, 2e+06, 3e+06, 4e+06, 5e+06, 6e+06, 7e+06
```

```
## $ run0 <dbl> 0.09924603, 0.19706607, 0.30280304, 0.44548678
```

```
## $ run1 <dbl> 0.1099961, 0.1945050, 0.3008888, 0.5314040, 0.7514040
```

```
## $ run2 <dbl> 0.1018548, 0.2033260, 0.3165970, 0.4161611, 0.5161611
```

```
## $ run3 <dbl> 0.1004527, 0.1933441, 0.3653409, 0.4889781, 0.5889781
```

```
## $ run4 <dbl> 0.1140921, 0.2565329, 0.3853610, 0.4850140, 0.5850140
```

Exercise 3

```
d <- gather(d, key="run_number", value="run_time",  
            run0, run1, run2, run3, run4)
```

Exercise 4

```
separate(d, run_number,  
         into=c("dummy", "run_number"), sep=3)
```

```
## # A tibble: 50 x 4  
##           n dummy run_number run_time  
##       <dbl> <chr> <chr>      <dbl>  
## 1 1000000 run      0      0.0992  
## 2 2000000 run      0      0.197  
## 3 3000000 run      0      0.303  
## 4 4000000 run      0      0.445  
## 5 5000000 run      0      0.584  
## 6 6000000 run      0      0.771  
## 7 7000000 run      0      1.54  
## 8 8000000 run      0      0.982  
## 9 9000000 run      0      1.24  
## 10 10000000 run      0      1.38  
## # ... with 40 more rows
```

Exercise 5

```
separate(d, run_number,  
         into=c(NA, "run_number"), sep=3)
```

```
## # A tibble: 50 x 3  
##           n run_number run_time  
##       <dbl> <chr>      <dbl>  
## 1 1000000 0          0.0992  
## 2 2000000 0          0.197  
## 3 3000000 0          0.303  
## 4 4000000 0          0.445  
## 5 5000000 0          0.584  
## 6 6000000 0          0.771  
## 7 7000000 0          1.54  
## 8 8000000 0          0.982  
## 9 9000000 0          1.24  
## 10 10000000 0         1.38  
## # ... with 40 more rows
```

Exercise 6

```
d <- separate(d, run_number,  
              into=c(NA, "run_number"), sep=3,  
              convert=TRUE)
```

```
d
```

```
## # A tibble: 50 x 3
```

```
##           n run_number run_time
```

```
##           <dbl>       <int>    <dbl>
```

```
## 1 10000000           0 0.0992
```

```
## 2 20000000           0 0.197
```

```
## 3 30000000           0 0.303
```

```
## 4 40000000           0 0.445
```

```
## 5 50000000           0 0.584
```

```
## 6 60000000           0 0.771
```

```
## 7 70000000           0 1.54
```

```
## 8 80000000           0 0.982
```

```
## 9 90000000           0 1.24
```

```
## 10 100000000         0 1.28
```

Exercise 7

```
write_csv(d, "data/sort_times_tidy.csv")
```

Exercises (dplyr)

- Exercise 1: Our sort times data is available in tidy format as `sort_times_tidy.csv`. Use `group_by` and `summarise` to get the mean and the standard deviation for each `n`, and then for each `run_number`.
- A dataset of characters in *Star Wars* is available as `dplyr::starwars`.
Exercise 2: Find all the human females. Exercise 3: Find the characters who are human *or* Wookiee. Exercise 4: Find the shortest character. Hint: recall we might need `na.rm`. Exercise 5: Add a new column called BMI giving the body mass index, where the formula is $BMI = m/h^2$ for mass m in kg and height h in metres. https://en.wikipedia.org/wiki/Body_mass_index. Exercise 6: Which character has the highest BMI?

Solutions (dplyr)

Exercise 1

```
d <- read_csv("data/sort_times_tidy.csv")
```

```
## Parsed with column specification:
```

```
## cols(
```

```
##   n = col_double(),
```

```
##   run_number = col_double(),
```

```
##   run_time = col_double()
```

```
## )
```

```
d %>% group_by(n) %>%
```

```
  summarise(mean_run_time=mean(run_time),
```

```
            sd_run_time=sd(run_time))
```

```
## # A tibble: 10 x 3
```

```
##           n mean_run_time sd_run_time
```

```
##           <dbl>         <dbl>         <dbl>
```

```
##  1  1000000          0.105         0.00654
```

```
##  2  2000000          0.209         0.0269
```

```
##  3  3000000          0.334         0.0387
```

Exercise 1

Notice that the mean and stddev for $n = 7$ million are anomalously high. One way this could occur is if our computer had a spike in CPU usage during the experiment, e.g. due to a browser loading a video.

Exercise 1

```
d %>% group_by(run_number) %>%  
  summarise(mean_run_time=mean(run_time),  
            sd_run_time=sd(run_time))
```

```
## # A tibble: 5 x 3
```

```
##   run_number mean_run_time sd_run_time  
##      <dbl>         <dbl>         <dbl>  
## 1         0         0.754         0.512  
## 2         1         0.644         0.368  
## 3         2         0.604         0.353  
## 4         3         0.648         0.369  
## 5         4         0.678         0.416
```

No major anomalies this time.

Exercise 2

```
sw <- dplyr::starwars
# human females
sw %>% filter(species == "Human", gender == "female")

## # A tibble: 9 x 13
##   name    height    mass hair_color skin_color eye_color birth_
##   <chr>    <int> <dbl> <chr>      <chr>      <chr>      <chr>
## 1 Leia~    150     49 brown      light      brown
## 2 Beru~    165     75 brown      light      blue
## 3 Mon ~    150     NA auburn     fair       blue
## 4 Shmi~    163     NA black      fair       brown
## 5 Cordé    157     NA brown     light      brown
## 6 Dormé    165     NA brown     light      brown
## 7 Joca~    167     NA white     fair       blue
## 8 Rey      NA      NA brown     light      hazel
## 9 Padm~    165     45 brown     light      brown
## # ... with 5 more variables: homeworld <chr>, species <chr>
```

Exercise 3

```
# human or Wookiee
```

```
sw %>% filter(species == "Human" | species == "Wookiee")
```

```
## # A tibble: 37 x 13
```

```
##   name   height  mass hair_color skin_color eye_color birth
```

```
##   <chr>   <int> <dbl> <chr>         <chr>      <chr>
```

```
## 1 Luke~    172    77 blond         fair       blue
```

```
## 2 Dart~    202   136 none          white      yellow
```

```
## 3 Leia~    150    49 brown         light      brown
```

```
## 4 Owen~    178   120 brown, gr~    light      blue
```

```
## 5 Beru~    165    75 brown         light      blue
```

```
## 6 Bigg~    183    84 black         light      brown
```

```
## 7 Obi~-    182    77 auburn, w~   fair       blue-gray
```

```
## 8 Anak~    188    84 blond         fair       blue
```

```
## 9 Wilh~    180    NA auburn, g~   fair       blue
```

```
## 10 Chew~   228   112 brown         unknown    blue
```

```
## # ... with 27 more rows, and 5 more variables: homeworld <
```

Exercise 4

```
sw %>% filter(height == max(height, na.rm=TRUE))

## # A tibble: 1 x 13
##   name height mass hair_color skin_color eye_color birth_
##   <chr>  <int> <dbl> <chr>      <chr>      <chr>      <chr>
## 1 Yara~    264   NA none        white        yellow
## # ... with 5 more variables: homeworld <chr>, species <chr>
## #   vehicles <list>, starships <list>
```

Exercise 5

```
# NB convert height from cm to metres before squaring  
BMI <- function(h, m) {m / (h / 100)^2}  
sw <- sw %>% mutate(bmi=BMI(height, mass))
```


Exercise 6

```
sw %>% filter(bmi == max(bmi, na.rm=TRUE))

## # A tibble: 1 x 14
##   name height mass hair_color skin_color eye_color birth_
##   <chr>  <int> <dbl> <chr>      <chr>      <chr>      <chr>
## 1 Jabb~    175  1358 <NA>      green-tan~ orange
## # ... with 6 more variables: homeworld <chr>, species <chr>
## #   vehicles <list>, starships <list>, bmi <dbl>
```