

dplyr joins

James McDermott

University of Galway



OLLSCOIL NA GAILLIMHE
UNIVERSITY OF GALWAY



Programming and Tools for Artificial Intelligence

Dr James McDermott

`dplyr joins`

Relational databases

The main ideas of relational databases (SQL) are probably familiar to all:

- A database consists of tables
- A table consists of a set of columns
- A column has a type, and maybe some constraints (e.g. positive integer)
- Some column(s) may be designated as a key for the table

- As we know, in Relational Databases, it is good practice to use *normalisation*: splitting a table up into multiple tables, to avoid duplication of information and the possibility of *update anomalies*. 3NF is the result of normalisation.
- Doing ML/stats/analytics may require *de-normalisation* – re-joining – eventually to export to our ML/stats/analytics system.

Before normalisation

Movie rental DB				
Date	Movie	Genre	Customer	Address
01-Jan	Amelie	Romance	Bob	11, Haight St
02-Jan	The Matrix	Sci-fi	Frida	Oxford Circus
02-Jan	Amelie	Romance	Carrie	99, Fifth Ave
05-Jan	Skyfall	Adventure	Bob	11, Haight St
05-Jan	Avengers	Sci-fi	Frida	Oxford Circus

After normalisation: 3rd Normal Form (3NF)

Rentals table		
Date	Movie ID	Customer ID
01-Jan	102	1
02-Jan	101	2
02-Jan	102	3
05-Jan	103	1
05-Jan	104	2
Customer table		
Customer ID	Name	Address
1	Bob	11, Haight St
2	Frida	Oxford Circus
3	Carrie	99, Fifth Ave
Movie table		
Movie ID	Name	Genre
101	The Matrix	Sci-fi
102	Amelie	Romance
103	Skyfall	Adventure
104	Avengers	Sci-fi

Key columns

After normalisation, the link between data is via key columns – in this case, the Customer ID and Movie ID columns. It is possible to put the original table back together using a **join**. We say that we join **on** the key column.

In SQL, a JOIN might be something like this. This is an *implicit join*:

```
SELECT * FROM RENTALS, CUSTOMER  
WHERE RENTALS.CustomerID = CUSTOMER.CustomerID;
```

This is an equivalent *explicit join*:

```
SELECT * FROM RENTALS JOIN CUSTOMER  
ON RENTALS.CustomerID = CUSTOMER.CustomerID;
```





(This is not examinable.)

What is a join, really?

- Think of join as an *operator* whose left and right operands are tables, and whose result is a table formed as the union of their columns
- The *Cross join* is a good place to start. Conceptually, a cross join is a *Cartesian product of rows*. For every row in T1, we put it side by side with every row in T2. Think of that as a new joined table. Now we can select columns from it and filter rows using ON. In particular, we'll probably filter for rows where a key column in one table matches a key column in the other, discarding the large majority of this cross product.
- Other joins just restrict the “every row in T1” and “every row in T2” parts depending on which matches actually exist.

Different types of joins

There are a few types of joins. To distinguish them, many textbooks and cheatsheets proceed to Venn diagrams, e.g. <http://www.sql-join.com/sql-join-types> (below). These are helpful as mnemonics but the language of Venn diagrams is not sufficient to define the different joins.



Different types of joins (Data Wrangling Cheatsheet)

a		b	
x1	x2	x1	x3
A	1	A	T
B	2	B	F
C	3	D	T

+

=

Mutating Joins

x1	x2	x3
A	1	T
B	2	F
C	3	NA

dplyr::left_join(a, b, by = "x1")
Join matching rows from b to a.

x1	x3	x2
A	T	1
B	F	2
D	T	NA

dplyr::right_join(a, b, by = "x1")
Join matching rows from a to b.

x1	x2	x3
A	1	T
B	2	F

dplyr::inner_join(a, b, by = "x1")
Join data. Retain only rows in both sets.

x1	x2	x3
A	1	T
B	2	F
C	3	NA
D	NA	T

dplyr::full_join(a, b, by = "x1")
Join data. Retain all values, all rows.

Filtering Joins

x1	x2
A	1
B	2

dplyr::semi_join(a, b, by = "x1")
All rows in a that have a match in b.

x1	x2
C	3

dplyr::anti_join(a, b, by = "x1")
All rows in a that do not have a match in b.

Further reading

- Most people working in industry in the fields of AI, ML, Data Science, Statistics, etc., use relational databases and SQL a lot.
- We don't teach it, because it is usually seen as a topic for undergrad level. This MOOC is recommended as an optional catch-up or refresher:
 - Stanford Databases
<https://lagunita.stanford.edu/courses/DB/2014/SelfPaced/about>
 - (The following topics in the MOOC are recommended for a “short version”: Introduction, JSON, Relational Algebra (Section 1), SQL, Relational Design Theory (Section 1), Unified Modelling Language, Online Analytical Processing)

Exercises

- 1 Read the three data files `rentals.csv`, `movies.csv`, `customers.csv`, all in the `data/` directory, as tibbles.
- 2 Optional: get R to read the `Date` column correctly. Hint: https://readr.tidyverse.org/reference/parse_datetime.html
- 3 Using a `dplyr` join command, create a table showing the customer name and address for every rental.
- 4 Piping the result into another join command, recreate the full original table as shown under “Before Normalisation” above.
- 5 Notice the columns `Name.x` and `Name.y` which appear because there is a `Name` column in each of the `Movies` and `Customers` tables. Rename them.
- 6 Calculate the number of movies Frida watched of the Sci-fi genre.

Solutions

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse
## v ggplot2 3.2.1      v purrr  0.3.2
## v tibble  2.1.3      v dplyr  0.8.3
## v tidyr   1.0.0      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0
## -- Conflicts ----- tidyverse_confli
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

Exercises 1 and 2:

```
rentals <- read_csv("data/rentals.csv",  
                    col_types=cols(Date=col_date(  
                      format="%d-%b-%Y")))  
movies <- read_csv("data/movies.csv")  
  
## Parsed with column specification:  
## cols(  
##   MovieID = col_double(),  
##   Name = col_character(),  
##   Genre = col_character()  
## )
```

```
customers <- read_csv("data/customers.csv")  
  
## Parsed with column specification:  
## cols(  
##   CustomerID = col_double(),  
##   Name = col_character(),  
##   Address = col_character()
```

Customer name and address for each rental

```
inner_join(rentals, customers, by="CustomerID")
```

```
## # A tibble: 5 x 5
```

##	Date	MovieID	CustomerID	Name	Address
##	<date>	<dbl>	<dbl>	<chr>	<chr>
## 1	2018-01-01	102	1	Bob	11, Haight St
## 2	2018-01-02	101	2	Frida	Oxford Circus
## 3	2018-01-02	102	3	Carrie	99, Fifth Ave
## 4	2018-01-05	103	1	Bob	11, Haight St
## 5	2018-01-05	104	2	Frida	Oxford Circus

Recreate original table

```
inner_join(rentals, customers, by="CustomerID") %>%  
  inner_join(movies, by="MovieID")
```

```
## # A tibble: 5 x 7
```

##	Date	MovieID	CustomerID	Name.x	Address	Name.y
##	<date>	<dbl>	<dbl>	<chr>	<chr>	<chr>
## 1	2018-01-01	102	1	Bob	11, Haight St	Amelie
## 2	2018-01-02	101	2	Frida	Oxford Circus	The Ma
## 3	2018-01-02	102	3	Carrie	99, Fifth Ave	Amelie
## 4	2018-01-05	103	1	Bob	11, Haight St	Skyfal
## 5	2018-01-05	104	2	Frida	Oxford Circus	Avenge

Rename columns

```
t = inner_join(rentals, customers, by="CustomerID") %>%  
  inner_join(movies, by="MovieID") %>%  
  rename(CustomerName=Name.x, MovieTitle=Name.y)
```

t

```
## # A tibble: 5 x 7
```

##	Date	MovieID	CustomerID	CustomerName	Address	Mo
##	<date>	<dbl>	<dbl>	<chr>	<chr>	<c
## 1	2018-01-01	102	1	Bob	11, Haight~	An
## 2	2018-01-02	101	2	Frida	Oxford Cir~	Th
## 3	2018-01-02	102	3	Carrie	99, Fifth ~	An
## 4	2018-01-05	103	1	Bob	11, Haight~	Sk
## 5	2018-01-05	104	2	Frida	Oxford Cir~	Av

Filter and count

```
t %>% filter(CustomerName=="Frida", Genre=="Sci-fi") %>%  
  count()
```

```
## # A tibble: 1 x 1  
##       n  
##   <int>  
## 1     2
```

Filter and count

The following is a solution to the problem, but it requires the programmer to do all the work in their head. That's not scalable or flexible and it's error-prone, so don't do this.

```
# Frida is CustomerID 2
# Movies 101 and 104 are Sci-fi
rentals %>% filter(CustomerID == 2,
                  MovieID %in% c(101, 104)) %>%
  count()

## # A tibble: 1 x 1
##       n
##   <int>
## 1     2
```