

Lecture 01 – Introduction

Optimisation CT5141

James McDermott

NUI Galway



OÉ Gaillimh
NUI Galway

Overview

- 1 What is Optimisation?**
- 2 Spherical rocks and special cases**
- 3 A binary guessing game**
- 4 Lots of applications**

What is Optimisation?



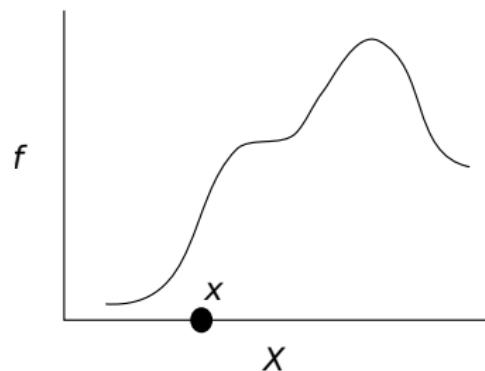
optimum: noun, from Latin, *optimus*, “best”

What is Optimisation?

“The Science of Better” <http://www.scienceofbetter.org/>

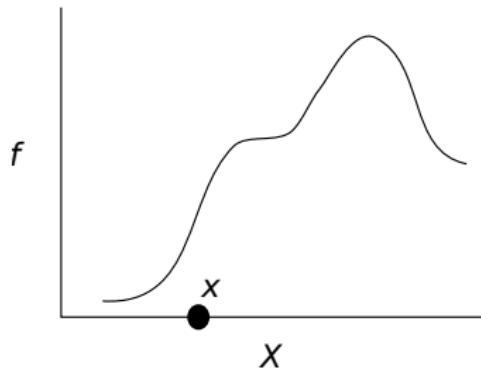
- It's about improving things, reducing costs, maximising profits, saving fuel, reducing wasted time, ...
- Optimisation methods are scientific/mathematical, but the process of formalising a problem can require creativity and intuition.

What is Optimisation?



We search for $x \in X$ which maximises the value of $f : X \rightarrow \mathbb{R}$

What is Optimisation?



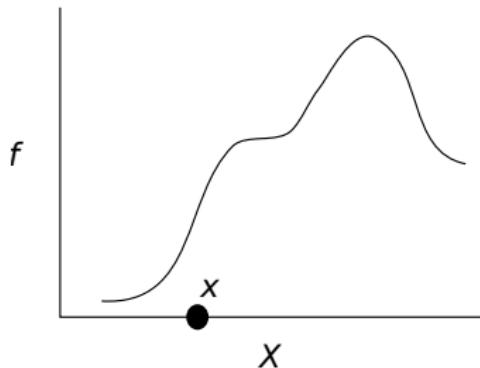
We search for $x \in X$ which maximises the value of $f : X \rightarrow \mathbb{R}$

- X is a **search space**, the set of possibilities.
- f is an **objective function**: $f(x) \in \mathbb{R}$ measures “how good” is x .
E.g., “how much energy is produced by the generator design x ?“

require
min

max

What is Optimisation?



We search for $x \in X$ which maximises the value of $f : X \rightarrow \mathbb{R}$

- X is a **search space**, the set of possibilities.
- f is an **objective function**: $f(x) \in \mathbb{R}$ measures “how good” is x .
E.g., “how much energy is produced by the generator design x ?”

In other words, we try to find: $\arg \max_{x \in X} f(x)$

What is Optimisation?

Optimisation = Search

What is not Optimisation?

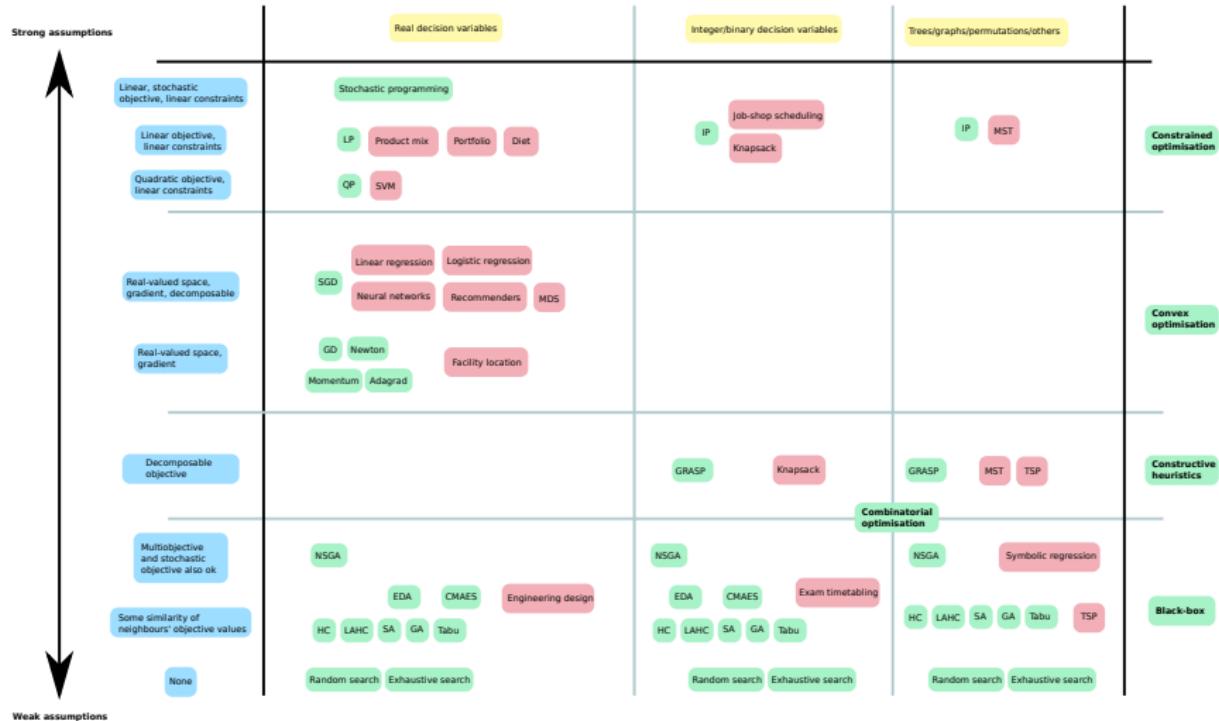
(... for our purposes)

- Rewriting an algorithm to achieve $O(\log n)$ instead of $O(n^2)$
- Optimising compilers
- Search in the sense of *pathfinding*, e.g. A* or breadth-first search.

Many problems, many algorithms

There are many optimisation problems and many algorithms. The right algorithm may depend on:

- The type of search space (e.g. real vectors, trees, permutations, bitstrings...)
- The properties of the objective (e.g. smooth, has a gradient, convex, multiobjective, ...)
- Constraints (i.e. some solutions are **invalid**)



Overview

- 1 What is Optimisation?
- 2 **Spherical rocks and special cases**
- 3 A binary guessing game
- 4 Lots of applications

Example: finding spherical rocks



How come there are lots of jagged rocks, irregular rocks, oval rocks – but no spherical rocks?

Example: finding spherical rocks



Suppose I wanted to **find** a spherical rock. Suppose I have a robot which can drive around, pick up rocks and scan them with a camera.

How could I use it to find the **most spherical** rock? I want it to do all this automatically, while I have lunch.

Example: finding spherical rocks

“Scanning” means it produces a data file that describes the 3D shape in detail. If you like, it’s just a cubic 3D array containing a 1 for every filled voxel and a 0 for every empty voxel.

E.g., this is a very small, flat, triangular rock:

```
[[[0,1,0],  
 [0,1,1],  
 [0,0,0]],  
 [[0,1,0],  
 [0,1,1],  
 [0,0,0]],  
 [[0,0,0],  
 [0,0,0],  
 [0,0,0]]]
```

Example: finding spherical rocks

Suppose we have a function that calculates `sphericalness(x)` for any given rock `x` in this 3D format. We can then just program the robot to try every single rock in the collection `X`, as follows:

```
best = None
bestf = -1
for x in X:
    if sphericalness(x) > bestf:
        best = x
        bestf = sphericalness(best)
```

Example: finding spherical rocks

Suppose we have a function that calculates `sphericalness(x)` for any given rock `x` in this 3D format. We can then just program the robot to try every single rock in the collection `X`, as follows:

```
best = None  
bestf = -1  
for x in X:  
    if sphericalness(x) > bestf:  
        best = x  
        bestf = sphericalness(best)
```

Python gives us a shorthand for that common idiom:

```
max(X, key=sphericalness)
```

Enumerative search

```
max(X) # X a list of numbers
```

Enumerative search

```
max(X) # X a list of numbers
```

```
max(X, key=f) # pass each x to f
```

Enumerative search

```
max(X) # X a list of numbers
```

```
max(X, key=f) # pass each x to f
```

This is $\arg \max_{x \in X} f(x)$. Enumerative search is **guaranteed** to find the **optimum** – the very best point in X – because it tries every point in X .

An objective function: sphericalness

So the only hard part is to define sphericalness. How would you do it?

Reflection

- If X is very large, what happens?
- Is there any way to get a (fairly good) result much faster?

When would we use enumerative search?

- If we need the optimum, not just a “very good” point
- If we have no “smarter” algorithm
- If time/computational budget allows
- What about computational complexity?

When would we use enumerative search?

- If we need the optimum, not just a “very good” point
- If we have no “smarter” algorithm
- If time/computational budget allows
- What about computational complexity?



It's a trap!

Random search

```
def rand_search(X, f, its=10000):
    return max(random.sample(X, its), key=f)
```

its is short for **iterations**.

Random search

```
def rand_search(X, f, its=10000):
    return max(random.sample(X, its), key=f)
```

its is short for **iterations**.

- Disadvantage: no guarantee of finding the optimum
- Advantage: can be as fast as we want

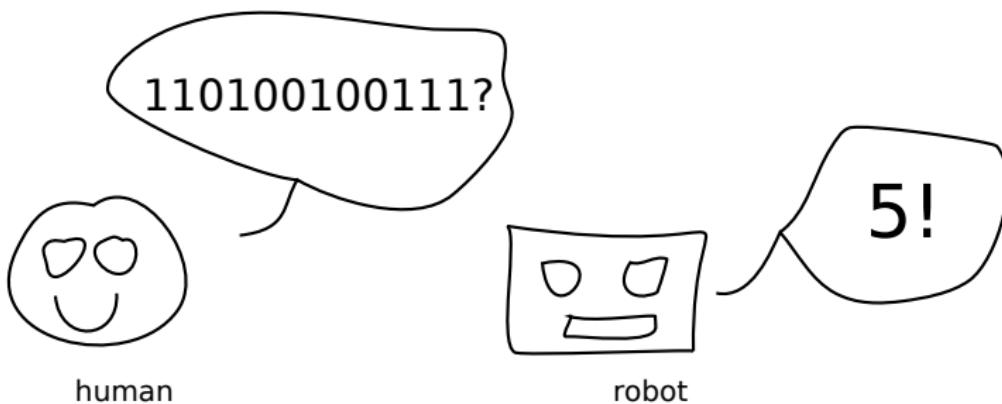
Reflection

- Random search is not very smart
- When we observe sphericalness (x_1) for a particular rock x_1 , what does it tell us about sphericalness (x_2) for the next rock x_2 ?

Overview

- 1 What is Optimisation?
- 2 Spherical rocks and special cases
- 3 A **binary guessing game**
- 4 Lots of applications

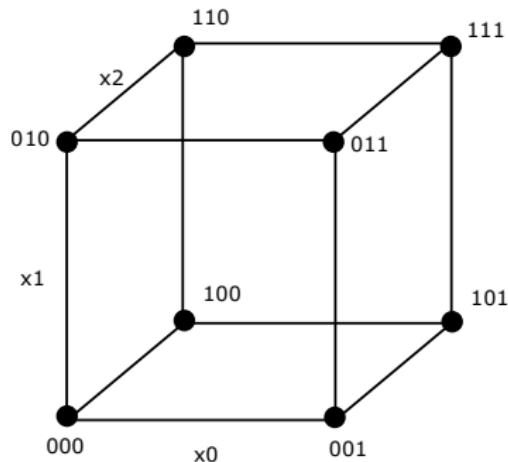
A guessing game (is an optimisation problem)



- A robot chooses a secret bitstring of length n
- We try to guess the bitstring
- Each time we guess, the robot tells us **how many bits are correct.**
- The bits are called **decision variables**.

What is the search space?

$X = \mathbb{H}^n$ (Hamming cube in n dimensions, i.e. space of bitstrings of length n)



What is the objective function?

Maximise: $f(x)$ = number of bits correct

What is the objective function?

Maximise: $f(x)$ = number of bits correct

Or equivalently, **minimise:** $f(x) = d_H(t, x)$, the **Hamming distance** between the guess x and the computer's target t .

What is the objective function?

Maximise: $f(x)$ = number of bits correct

Or equivalently, **minimise:** $f(x) = d_H(t, x)$, the **Hamming distance** between the guess x and the computer's target t .

secret	110100100111
guess	110000110111
<hr/>	
differences	000100010000
Hamming distance	2

Hill-climbing

```
def hill_climb(init, nbr, f, its=10000):
    x = init()
    for i in range(its):
        xnew = nbr(x) # try a new point...
        if f(xnew) > f(x):
            x = xnew # improvement!
    return x

def init(): # random bitstring
    return [random.randrange(2) for i in range(n)]
def nbr(x): # make a "neighbour" of x
    x = x.copy()
    i = random.randrange(n)
    x[i] = 1 - x[i] # flip a random bit
    return x
```

Hill-climbing

- Hill-climbing may be the simplest true optimisation algorithm
- Notice that the structure is similar to our `max` code: initialise x , then every time we find an improvement, replace x .
- We'll study this properly in a few weeks.

The curse of dimensionality

The curse of dimensionality: when the number of dimensions grows, every problem (machine learning, optimisation, search, ...) gets harder, and soon becomes intractable.

Reflection

- 1 In the guessing game, how does the size of the space grow as bitstring length n grows?

Reflection

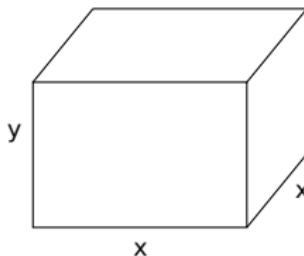
- 1 In the guessing game, how does the size of the space grow as bitstring length n grows?
- 2 For each guess x , we received $f(x)$, a single number.

What if, instead, we received more information – e.g. exactly **which** bits are incorrect?

Overview

- 1** What is Optimisation?
- 2** Spherical rocks and special cases
- 3** A binary guessing game
- 4** Lots of applications

Applications: Design



A box having a square base and an open top is to contain 108 cubic feet. What should its dimensions be to minimise materials? From
[http://www.themathpage.com/
aCalc/applied.htm](http://www.themathpage.com/aCalc/applied.htm).

We use “good old differentiation”: we formulate M , the amount of materials, as a function of x (the side of the base): $M = x^2 + 432/x$. We can just set $dM/dx = 0$ (and also check the second derivative).

Applications: Travelling salesman problem

X is the set of possible tours of some cities. f is a measure of **total tour distance**. We might use a **constructive heuristic**.



Applications: Energy production schedules



Given data on regional electricity requirements and power plants, design a schedule to switch plants on/off to meet demand. X is all possible schedules. f_1 represents cost of fuel, and f_2 total emissions. (The “unit commitment” problem.) We might use **multiobjective optimisation**.

Applications: Timetabling

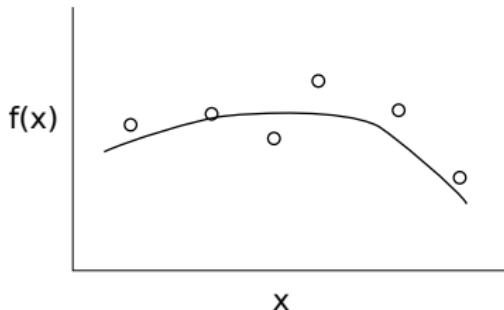
- X is the set of possible exam timetables. We have a **constraint** that any timetable with a clash is **invalid**.
- $f(x)$ might measure the number of students who have one exam right after another.
- We might use a **genetic algorithm**.

Applications: Portfolio allocation

Given a fixed amount of money to be invested and a set of n stocks to invest in, we want to maximise expected return and minimise risk.

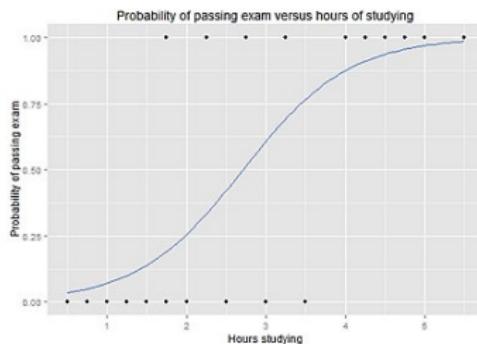
- $X \subset \mathbb{R}^n$ is the set of possible allocations. There are n decision variables.
- f is a measure of expected return and risk.
- We might use **linear programming**.

Applications: Curve-fitting



- Given a non-linear regression model $f(x)$ with fixed form but unknown parameters c
- e.g. $f(x) = c_0 x^{c_1} \cdot \log(c_2 x)$
- and given training data X and y
- we try to find the parameters c which minimise the loss $\sum_i(f(x_i) - y_i)^2$. These are the **decision variables** in our optimisation.
- We might use **non-linear least squares**.

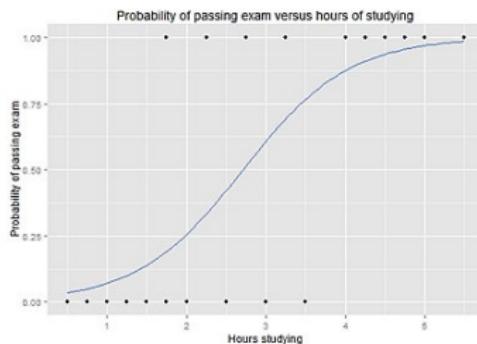
Applications: Machine Learning



Wiki

- Logistic regression: X is the set of possible **weight values** for the LR model. f might be the **loss** on training data. We might use **gradient descent**.

Applications: Machine Learning



Wiki

- Logistic regression: X is the set of possible **weight values** for the LR model. f might be the **loss** on training data. We might use **gradient descent**.
- If we have a neural network, we need **gradient descent with backpropagation**.

Applications: Program synthesis

- Given a set of input-output examples (X, y) , we try to find a program $p \in X$ such that $p(x_i) = y_i$. f might measure how many examples our program p gets correct. We might use **genetic programming**.

Applications: Cache expiry algorithms

- A size-limited hardware or software cache uses some algorithm to decide which items to throw out when it becomes full, e.g. a least-recently used (LRU) algorithm.
- X is the set of possible parameters of that algorithm and f could be a measure of throughput.
- We might use simulation to calculate f . We might use **covariance matrix adaptation** to search for x .

Applications: Menus



From topsecretrecipes.com

- Suppose I'm going to host a party for 10 small kids. They'll eat sausages (EUR5/kg), chips (EUR2/kg), and ice-cream (EUR4/kg). The kids don't care what they get so long as they get 500g of food each. Suppose I don't care about their health, only about cost. How much of each type of food should I buy?